

Mean Field Game Modeling of Overconfidence in Financial Markets

Lucas Sneller

Miami University, Oxford, OH 45056, USA

Abstract

This paper develops a mean field model of overconfidence in a market with flow-driven price impact and inventory carryover. A continuum of investors filters an unobserved fundamental from noisy private signals and trades against an endogenous price that mean-reverts toward fundamentals and responds to aggregate order flow; overconfidence enters as misperceived private-signal precision, generating heterogeneous filtered beliefs. Our main equilibrium object is a discrete-time *myopic mean-field equilibrium*: given a conjectured mean trading rate, each agent chooses next-step inventory to maximize one-step CARA utility under conditionally Gaussian price increments. This yields a linear feedback rule and reduces mean-field consistency to a scalar stock-flow fixed point at each time step. We provide an explicit per-step contraction condition guaranteeing existence, uniqueness, and stability of the resulting discrete-time dynamics and particle approximation. In simulations across bull/bear regimes and a targeted joint grid over anchoring, impact, and noise trading, microstructure parameters govern mispricing levels, whereas overconfidence primarily increases disagreement and trading intensity with limited effect on price-level mispricing in the baseline anchored-impact regime. A finite-horizon continuous-time CARA/LQG solution is reported as a benchmark, and an identification/moment-matching vignette connects the model's parameters to price and volume moments.

1 Introduction

Overconfidence—the tendency to overestimate the precision of one's information—is a well-documented behavioral bias in financial markets and is empirically associated with excessive trading and episodes of mispricing Odean (1998); Barber and Odean (2001). A central modeling challenge is to combine heterogeneous (and potentially biased) beliefs formed from private information with endogenous price formation, while retaining an equilibrium concept that is mathematically well-posed, computable, and connectable to data. This paper develops a mean field game (MFG) framework for a market populated by a continuum of investors who filter an unobserved fundamental from noisy private signals and trade against an endogenous price impacted by mean trading rate (order flow).

The model keeps a clean stock-flow separation. Inventory $x_{i,t}$ (shares) carries over between decision times, trading rate $u_{i,t}$ (shares per unit time) satisfies $x_{i,t+\Delta t} = x_{i,t} + u_{i,t}\Delta t$, and the price equation loads on the mean flow $\bar{u}_t$ rather than on mean inventory. Overconfidence enters only through *filtering*: agents treat their private signal as more precise than it truly is, which changes posterior uncertainty and therefore trading intensity and disagreement, not through an ad hoc demand multiplier.

Early behavioral models introduced stylized agent types (noise traders, momentum traders) to demonstrate how bounded rationality can drive prices away from fundamentals De Long et al. (1990); Hong and Stein (1999); Daniel et al. (1998); Gervais and Odean (2001), but typically

rely on a small number of types rather than a full distribution of beliefs. Mean field games have been applied to financial markets in several contexts: Guéant, Lasry, and Lions Guéant et al. (2010) developed MFG frameworks for execution and portfolio problems; Carmona, Fouque, and Sun Carmona et al. (2015) studied systemic risk; Lachapelle *et al.* Lachapelle et al. (2016) applied MFG ideas to high-frequency trading and price formation; and Cardaliaguet and Lehalle Cardaliaguet and Lehalle (2018) analyzed trade crowding effects.

Our contribution relative to existing MFG finance literature is to provide a bounded-rational equilibrium concept that is both economically interpretable and mathematically checkable in a stock–flow microstructure setting. In the dynamic baseline, we study a discrete-time *myopic mean-field equilibrium*: given a conjectured common-noise adapted mean trading rate ($\bar{u}_t$), each agent applies a one-step CARA best response under conditionally Gaussian price increments (using their filtered belief), and mean-field consistency closes the dynamics by requiring $\bar{u}_t$ to equal the conditional population mean of these best responses. Because the best response is linear in $(\hat{v}_{i,t} - p_t)$ and in $\bar{u}_t$, the closure reduces each time step to a scalar stock–flow fixed point. A per-step contraction condition guarantees uniqueness and well-posedness of the induced discrete-time dynamics and particle approximation (Theorem 4.6).

The simulations distinguish *level amplification* (large mispricing/volatility levels driven by weak anchoring and high noise trading) from *overconfidence effects* (incremental differences when increasing k). Across a targeted joint grid over anchoring, impact, and noise trading, we find a robust mechanism separation in the baseline anchored-impact regime: microstructure parameters govern mispricing levels, while overconfidence strongly increases disagreement and trading intensity with economically small effects on price-level mispricing (Subsection 6.2.1). To make the model empirically usable, Section 6.1 provides an identification argument linking key parameters to observable price and volume moments, and a moment-match vignette illustrates the calibration bridge on a liquid ETF.

The paper’s dynamic analysis has a deliberate two-model structure. (A) The object we compute and simulate is the discrete-time stock–flow myopic equilibrium described above. (B) As a theoretical benchmark, Subsection 4.3 derives an exact finite-horizon continuous-time CARA/LQG best response under partial information; this analytic extension clarifies how the full intertemporal solution adds a hedging correction to the myopic rule, and we compare the myopic and intertemporal policies in Subsection 6.3.

The main contributions of this paper are:

- We formulate a mean field game with overconfident Bayesian filtering and an explicit stochastic price equation with endogenous flow impact, under a stock–flow timing that separates inventory and trading rate (Subsubsection 2.3.3).
- We provide a clean static benchmark with a closed-form linear price (5) and explicit comparative statics in the overconfidence parameter k .
- In the dynamic setting, we derive a checkable myopic CARA inventory rule under conditionally Gaussian price increments (with holding cost) and obtain a per-step stock–flow mean-field closure; an explicit contraction condition guarantees uniqueness and discrete-time well-posedness (Theorem 4.6).
- We report particle-based numerical experiments (large- N simulations) across bull/bear regimes and microstructure regimes, including mechanism diagnostics (belief dispersion and perceived uncertainty), validation checks, and extensions (Section 6).
- We provide a minimal identification/calibration bridge: Section 6.1 links key parameters to observable price and volume moments and illustrates the mapping with a moment-match vignette.

The remainder of this paper is organized as follows. Section 2 presents the model, including the static benchmark and the dynamic price/belief dynamics. Section 3 discusses the mean-field object in the presence of common noise and the fixed-point notion of (myopic) equilibrium. Section 4 states the main analytic results, including the myopic rule and the continuous-time CARA/LQG benchmark. Section 5 collects proofs. Section 6 presents numerical experiments, the empirical bridge, and the joint-grid analysis. Section 7 concludes.

Main contributions

- We formulate a mean field game with overconfident Bayesian filtering and an explicit stochastic price equation with endogenous flow impact, under a stock–flow timing that separates inventory and trading rate (Subsubsection 2.3.3).
- We provide a clean static benchmark with a closed-form linear price (5) and explicit comparative statics in the overconfidence parameter k .
- In the dynamic setting, we derive a checkable myopic CARA inventory rule under conditionally Gaussian price increments (with holding cost) and obtain a per-step stock–flow mean-field closure; an explicit contraction condition guarantees uniqueness and discrete-time well-posedness (Theorem 4.6).
- We report particle-based numerical experiments (large- N simulations) across bull/bear regimes and microstructure regimes, including mechanism diagnostics (belief dispersion and perceived uncertainty), validation checks, and extensions (Section 6).
- We provide a minimal identification/calibration bridge: Section 6.1 links key parameters to observable price and volume moments and illustrates the mapping with a moment-match vignette.

2 Problem Statement and Model

2.1 Setup and Notation

We consider a financial market with a continuum of investors indexed by $i \in [0, 1]$ over a horizon $[0, T]$. Let $(\Omega, \mathcal{F}, \{\mathcal{F}_t\}_{t \geq 0}, \mathbb{P})$ support independent Brownian motions W_t (fundamental shock), Z_t (aggregate price/noise-trading shock), U_t (common signal component), and $\{B_{i,t}\}_{i \in [0,1]}$ (idiosyncratic signal noise). We write $\mathbb{E}[\cdot]$ for expectation and $\tau = 1/\sigma^2$ for precision.

Model parameters: $\gamma > 0$ (absolute risk aversion), $\kappa > 0$ (fundamental anchoring strength), $\lambda > 0$ (impact strength), $k \geq 1$ (overconfidence), and volatilities $\sigma_v, \sigma_c, \sigma_\epsilon, \sigma_\eta, \phi > 0$ (fundamental shock, common signal component, idiosyncratic signal noise, order-flow noise scale, and reduced-form price/noise-trading volatility). In the microstructure sketch below, σ_η induces price noise via $\phi := \lambda \sigma_\eta$. Overconfidence is modeled as a misperception of private-signal precision: an agent with parameter k uses perceived idiosyncratic variance σ_ϵ^2/k in belief updating (equivalently, perceived precision k/σ_ϵ^2).

2.2 Static Benchmark

In this static benchmark, $\nu \sim \mathcal{N}(0, \sigma_\nu^2)$ is an exogenous noise-demand shock (price level), distinct from the order-flow noise scale σ_η used in the microstructure sketch (whose induced price-noise volatility is $\phi = \lambda \sigma_\eta$).

As a baseline, consider a static market with fundamental value $v \sim \mathcal{N}(\mu_{v,0}, \sigma_{v,0}^2)$. Each investor receives a private signal

$$s_i = v + \epsilon_i, \quad \epsilon_i \sim \mathcal{N}(0, \sigma_{\epsilon,0}^2), \quad (1)$$

and chooses demand x_i for a single risky asset with payoff v at settlement. Investors have CARA utility

$$U_i = -\mathbb{E}[e^{-\gamma W_i} \mid s_i], \quad W_i = w_0 + x_i(v - p), \quad (2)$$

and take the price impact rule (instantaneous clearing with noise demand)

$$p = \lambda \int_0^1 x_i di + \nu, \quad \nu \sim \mathcal{N}(0, \sigma_\nu^2). \quad (3)$$

Under Gaussian beliefs, maximizing CARA utility is equivalent to maximizing mean minus $(\gamma/2)$ times variance. Let $\tau_{v,0} := 1/\sigma_{v,0}^2$ and let $\tau'_{\epsilon,0} := k/\sigma_{\epsilon,0}^2$ denote the perceived private-signal precision. Then the (perceived) posterior variance is $\sigma_{\text{post}}^2 = 1/(\tau_{v,0} + \tau'_{\epsilon,0})$ and the posterior mean is $\hat{v}_i = (\tau_{v,0}\mu_{v,0} + \tau'_{\epsilon,0}s_i)/(\tau_{v,0} + \tau'_{\epsilon,0})$. The optimal demand is therefore

$$x_i^* = \frac{\hat{v}_i - p}{\gamma \sigma_{\text{post}}^2}. \quad (4)$$

To make the aggregation step explicit, consider first a finite population of size N , with signals $s_i = v + \epsilon_i$ for $i = 1, \dots, N$, where $\{\epsilon_i\}_{i=1}^N$ are i.i.d., independent of v , with $\mathbb{E}[\epsilon_i] = 0$ and $\mathbb{E}[\epsilon_i^2] = \sigma_{\epsilon,0}^2 < \infty$, and define the sample average $\bar{s}_N := \frac{1}{N} \sum_{i=1}^N s_i$. Since $\hat{v}_i$ is affine in s_i , the cross-sectional mean belief satisfies $\frac{1}{N} \sum_{i=1}^N \hat{v}_i = \frac{\tau_{v,0}\mu_{v,0} + \tau'_{\epsilon,0}\bar{s}_N}{\tau_{v,0} + \tau'_{\epsilon,0}}$. Under the finite- N clearing rule $p^N = \lambda(\frac{1}{N} \sum_{i=1}^N x_i) + \nu$, the same algebra yields

$$p^N = \frac{\lambda \tau'_{\epsilon,0}}{\gamma + \lambda(\tau_{v,0} + \tau'_{\epsilon,0})} \bar{s}_N + \frac{\gamma}{\gamma + \lambda(\tau_{v,0} + \tau'_{\epsilon,0})} \nu + \frac{\lambda \tau_{v,0}}{\gamma + \lambda(\tau_{v,0} + \tau'_{\epsilon,0})} \mu_{v,0}.$$

Since $\bar{s}_N = v + \frac{1}{N} \sum_{i=1}^N \epsilon_i$, the strong law of large numbers implies $\bar{s}_N \rightarrow v$ almost surely (conditional on v , hence also unconditionally) as $N \rightarrow \infty$, yielding the mean-field limit price¹

$$\begin{aligned} p &= \frac{\lambda \tau'_{\epsilon,0}}{\gamma + \lambda(\tau_{v,0} + \tau'_{\epsilon,0})} v + \frac{\gamma}{\gamma + \lambda(\tau_{v,0} + \tau'_{\epsilon,0})} \nu \\ &\quad + \frac{\lambda \tau_{v,0}}{\gamma + \lambda(\tau_{v,0} + \tau'_{\epsilon,0})} \mu_{v,0}. \end{aligned} \quad (5)$$

Equation (5) makes the channel explicit: higher k increases perceived precision $\tau'_{\epsilon,0}$ and therefore strengthens the price response to v while also increasing demand magnitude for a given mispricing.

Notation note. In the static benchmark (Sec. 2.2), $(\mu_{v,0}, \sigma_{v,0})$ parameterize the one-shot prior $v \sim \mathcal{N}(\mu_{v,0}, \sigma_{v,0}^2)$ and $\sigma_{\epsilon,0}$ is the one-shot private-signal noise standard deviation in $s_i = v + \epsilon_i$. We retain the “0” subscripts throughout Section 2.2 and its proofs; unsubscripted $(\mu_v, \sigma_v, \sigma_\epsilon, \dots)$ are reserved for the dynamic model. In the dynamic model (Sec. 2.3), (μ_v, σ_v) denote the drift and instantaneous diffusion volatility of v_t in $dv_t = \mu_v dt + \sigma_v dW_t$, and σ_ϵ is the diffusion loading on idiosyncratic signal noise in $d\xi_{i,t} = v_t dt + \sigma_c dU_t + \sigma_\epsilon dB_{i,t}$. The benchmarks are comparable but the parameters have different time-aggregation interpretations.

¹A literal continuum model with $\int_0^1 s_i di = v$ can be formalized under an “exact law of large numbers” framework for a continuum of independent shocks (e.g., via a Fubini extension/Loeb space). We avoid this measure-theoretic detour and adopt the large- N interpretation, consistent with the finite- N particle/limit perspective used later.

2.3 Dynamic Mean Field Game Formulation

2.3.1 State Dynamics and Mean-Field Coupling

The fundamental value follows a diffusion

$$dv_t = \mu_v dt + \sigma_v dW_t, \quad (6)$$

and each investor observes a private signal process

$$d\xi_{i,t} = v_t dt + \sigma_c dU_t + \sigma_\epsilon dB_{i,t}, \quad (7)$$

where $\{B_{i,t}\}$ are independent across i . In the discrete-time implementation on a grid with step size Δt , we work with the normalized signal increment (we reserve $y_{i,t}$ below for mispricing $\hat{v}_{i,t} - p_t$):

$$z_{i,t} := \frac{\xi_{i,t+\Delta t} - \xi_{i,t}}{\Delta t} = v_t + \frac{\sigma_c}{\sqrt{\Delta t}} \varepsilon_t^c + \frac{\sigma_\epsilon}{\sqrt{\Delta t}} \varepsilon_{i,t}^i, \quad \varepsilon_t^c, \varepsilon_{i,t}^i \sim \mathcal{N}(0, 1),$$

so $\text{Var}(z_{i,t} | v_t) = (\sigma_c^2 + \sigma_\epsilon^2)/\Delta t$ and the perceived measurement variance is $R'_i(\Delta t) = (\sigma_c^2 + \sigma_\epsilon^2/k_i)/\Delta t$ (Appendix A). (For $\Delta t = 1$, this reduces to $z_{i,t} = v_t + \sigma_c \varepsilon_t^c + \sigma_\epsilon \varepsilon_{i,t}^i$.) The market price evolves endogenously through an anchored impact rule:

$$dp_t = \kappa(v_t - p_t) dt + \lambda \bar{u}_t dt + \phi dZ_t, \quad \bar{u}_t := \int_0^1 u_{i,t} di. \quad (8)$$

Equation (8) is a reduced-form microstructure: prices mean-revert to fundamentals at rate κ (arbitrage/fundamental anchoring), respond to aggregate net trading flow with strength λ , and are perturbed by aggregate noise Z_t with volatility ϕ . The anchoring term $\kappa(v_t - p_t)$ should be interpreted as reduced-form arbitrage pressure toward the (latent) fundamental v_t ; it is not assumed that a market maker directly observes v_t . The mean-field coupling enters through $\bar{u}_t$. In the discrete-time implementation, $u_{i,t}$ is the trading rate induced by inventory adjustments over a decision interval Δt (Subsubsection 2.3.3).

Microstructure sketch for the anchored impact rule. Equation (8) can be viewed as a reduced-form limit of competitive risk-neutral market making with inventory costs, plus an (exogenous) arbitrage/anchoring flow.

Linear impact from inventory-averse competitive market makers. Let Q_t denote aggregate net order flow hitting dealers (positive when the public buys), and let the representative dealer's inventory be q_t , with

$$dq_t = -dQ_t, \quad (\text{dealers take the opposite side of net flow}).$$

Assume dealers are risk-neutral and competitive (zero expected profit), but face a quadratic inventory holding cost $\frac{\lambda}{2} q_t^2$. A standard myopic quoting / LQ argument implies the midprice is shifted linearly by inventory,

$$p_t \approx \tilde{m}_t - \lambda q_t,$$

where $\tilde{m}_t$ is the dealer's "efficient" benchmark (e.g., a fast-moving reference level). Differentiating and substituting $dq_t = -dQ_t$ gives

$$dp_t \approx d\tilde{m}_t + \lambda dQ_t,$$

so that (in drift form) price changes are *linear* in contemporaneous net flow. Writing $dQ_t = \bar{u}_t dt + \sigma_\eta dZ_t$ (where $\bar{u}_t$ is the net order-flow intensity hitting dealers) yields the impact and noise terms

$$dp_t - d\tilde{m}_t \approx \lambda \bar{u}_t dt + (\lambda \sigma_\eta) dZ_t, \quad \phi := \lambda \sigma_\eta.$$

Here $\tilde{m}_t$ is the dealer's efficient benchmark, $\bar{u}_t$ is the net order-flow intensity hitting dealers, and Z_t is a Brownian motion; σ_η is the corresponding order-flow noise scale. In the baseline below we explicitly distinguish inventory $x_{i,t}$ from trading rate $u_{i,t}$ (Subsubsection 2.3.3); impact in (8) depends on the mean trading rate $\bar{u}_t$ (equivalently $\lambda \Delta \bar{x}_t$ over a decision interval).

Anchoring as reduced-form arbitrage pressure. To capture fast arbitrage / fundamental-anchoring, posit an additional “arbitrage” flow that leans against mispricing,

$$u_t^{\text{arb}} = a(v_t - p_t),$$

so total net flow is $\bar{u}_t + u_t^{\text{arb}}$ at the same impact slope λ . Then

$$dp_t \approx \lambda(\bar{u}_t + a(v_t - p_t)) dt + (\lambda \sigma_\eta) dZ_t = \kappa(v_t - p_t) dt + \lambda \bar{u}_t dt + \phi dZ_t, \quad \kappa := \lambda a,$$

which matches the anchored impact specification.

2.3.2 Equilibrium Price Formation

We focus on feedback strategies where demand depends on perceived mispricing. Let $\hat{v}_{i,t}$ denote investor i 's Kalman–Bucy estimate of v_t under the *perceived* measurement-noise variance $R'(k)$ (Appendix A). When $k = 1$ this coincides with the true conditional mean $\mathbb{E}[v_t | \mathcal{F}_t^i]$; when $k \neq 1$ it should be interpreted as a *subjective* filter state computed under the investor's misperceived signal precision. Define the mean belief $m_t := \int_0^1 \hat{v}_{i,t} di$. A broad class of (best-response and learning-based) specifications take the form

$$x_{i,t} = \Theta_t(\Sigma_{i,t})(\hat{v}_{i,t} - p_t), \tag{9}$$

where $\Sigma_{i,t}$ is the agent's perceived posterior variance. *Definition (gain).* For each t , $\Theta_t : (0, \infty) \rightarrow [0, \infty)$ maps perceived posterior variance to the (shares/\$) sensitivity of demand to mispricing and is weakly decreasing in Σ (equivalently increasing in precision $1/\Sigma$). In the myopic CARAbest response, Θ_t is derived explicitly in Eq. (30) and is nonnegative under $\gamma > 0$, $\chi \geq 0$, and $\sigma_{p,i,t}^2(\Sigma) > 0$.

2.3.3 Timing and stock–flow separation (inventory vs. trading rate)

Timing and objects (discrete decision interval Δt). At each decision date t agents enter with a *pre-trade inventory* $x_{i,t-}$ (shares) and choose a *post-trade inventory* $x_{i,t}$ (shares). The executed order and trading rate are

$$\Delta x_{i,t} := x_{i,t} - x_{i,t-}, \quad u_{i,t} := \frac{\Delta x_{i,t}}{\Delta t} \quad (\text{shares/time}), \quad \bar{u}_t := \int_0^1 u_{i,t} di.$$

Price changes over $[t, t + \Delta t]$ via *flow-based* impact,

$$\Delta p_{t+\Delta t} = \kappa(v_t - p_t) \Delta t + \lambda \bar{u}_t \Delta t + \phi \sqrt{\Delta t} \varepsilon_{t+\Delta t}, \quad \varepsilon_{t+\Delta t} \sim \mathcal{N}(0, 1).$$

Equivalently, since $\bar{u}_t \Delta t = \Delta \bar{x}_t$ with $\bar{x}_t := \int_0^1 x_{i,t} di$, the impact term can be written $\lambda \Delta \bar{x}_t$. This timing makes explicit that x is a stock (shares) while u is a flow (shares/time), resolving the unit ambiguity when a single symbol is used for both.

2.3.4 Belief Updating via Kalman Filtering

Each investor maintains a belief state $(\hat{v}_{i,t}, \Sigma_{i,t})$ where $\hat{v}_{i,t}$ is the posterior mean of v_t and $\Sigma_{i,t}$ is the perceived posterior variance. In the linear-Gaussian setting above, $(\hat{v}_{i,t}, \Sigma_{i,t})$ satisfy Kalman–Bucy equations driven by the private signal $\xi_{i,t}$. Overconfidence enters by replacing σ_ϵ^2 with σ_ϵ^2/k in the perceived filter. We summarize the resulting filter in Appendix A.

2.3.5 Objective Functional

Investor i chooses a post-trade inventory process $(x_{i,t})_t$ (shares) on the discrete decision grid. Over one step, wealth evolves as

$$W_{i,t+\Delta t} = W_{i,t} + x_{i,t} \Delta p_{t+\Delta t} - \frac{\chi}{2} x_{i,t}^2 \Delta t, \quad \chi > 0,$$

where the quadratic term penalizes inventory holding/risk over the interval. Rather than solving the full intertemporal exponential-utility control problem, we adopt a bounded-rational *myopic* approximation: at each decision date, the investor solves the one-step certainty-equivalent problem

$$\max_{x_{i,t}} x_{i,t} \mu_{i,t} - \frac{1}{2} (\gamma \sigma_{p,i,t}^2 + \chi) x_{i,t}^2 \Delta t,$$

where $\mu_{i,t} := \mathbb{E}_i[\Delta p_{t+\Delta t} \mid \mathcal{F}_t^i] = (\kappa(\hat{v}_{i,t} - p_t) + \lambda \bar{u}_t) \Delta t$ and $\text{Var}_i(\Delta p_{t+\Delta t} \mid \mathcal{F}_t^i) = \sigma_{p,i,t}^2 \Delta t$. We treat $\sigma_{p,i,t}^2$ as computed from primitives; in the discrete implementation it includes belief uncertainty through the effective variance term in Section IV.

2.3.6 Mean-Field Consistency (Myopic)

Let μ_t denote the (possibly random) law of $(\hat{v}_{i,t}, \Sigma_{i,t})$ across agents. Together with (8), we study a *myopic mean-field consistent* fixed point: (i) given a conjectured mean-field flow $\{\mu_t\}$ (equivalently, moments such as m_t), each agent applies the myopic CARA rule in Section IV; and (ii) the resulting population law induced by the controlled belief dynamics is consistent with $\{\mu_t\}$. Section III formalizes the mean-field object and discusses the common-noise aspect induced by W_t and Z_t .

3 Mean-Field Approximation and Limiting Equations

3.1 Mean-Field Object Under Common Noise

Because the model includes common shocks (W_t, U_t, Z_t) through the fundamental, correlated signals, and price noise, the appropriate mean-field object is the *conditional* population law given the common filtration $\mathcal{F}_t^0 := \sigma(W_s, U_s, Z_s : s \leq t)$. Concretely, let $X_{i,t} := (\hat{v}_{i,t}, \Sigma_{i,t}, k_i)$ denote investor i 's internal state (belief, perceived uncertainty, and overconfidence type), and define the mean-field flow

$$\mu_t := \mathcal{L}(X_{i,t} \mid \mathcal{F}_t^0). \tag{10}$$

The conditional mean belief $m_t := \int \hat{v} \mu_t(d\hat{v} d\Sigma dk)$ enters the price dynamics through the mean trading rate $\bar{u}_t$ appearing in (8); inventories and trading rates are linked by the stock–flow timing in Subsubsection 2.3.3.

3.2 Mean-Field Equilibrium (Common-Noise Formulation)

Fix a candidate flow $\{\mu_t\}_{t \in [0, T]}$ and the resulting price process $\{p_t\}$ from (8). In the bounded-rational formulation of this paper, the representative agent does not solve a full intertemporal control problem; instead, given $\{\mu_t\}$ and $\{p_t\}$ the agent applies the myopic rule from Section IV, yielding a feedback $x_{i,t} = x^{\text{myopic}}(t, X_{i,t}, p_t, \mu_t)$. The induced controlled state process generates a new conditional flow $\{\tilde{\mu}_t\}$. A myopic mean-field equilibrium is a fixed point $\mu = \tilde{\mu}$.

In common-noise settings, μ_t is random and its evolution is governed by a stochastic Kolmogorov forward equation (stochastic Fokker–Planck / SPDE) for the conditional law; see, e.g., Carmona and Delarue (2018a); Lacker (2016). In Section VI we compute equilibrium statistics using a particle approximation: simulate a large N -agent system driven by shared (W, Z) and independent $\{B_i\}$ and estimate μ_t via the empirical measure.

Scope of theoretical results. To avoid confusion, we explicitly state what this paper establishes and what it does not claim. **Baseline object (what we compute and simulate):** a discrete-time, stock–flow myopic equilibrium with inventory carryover and flow-based impact. The per-step mean-field fixed point is scalar and has a unique solution when $0 \leq B_t < \Delta t$ (Theorem 4.6); under this condition the full discrete-time particle system is well-posed (Theorem 3.1). **Benchmarks (for comparison only):** we also record a continuous-time position-impact benchmark (which collapses stock and flow) where the mean-field closure becomes algebraic and admits an explicit closed form and condition (Theorem 3.2); this benchmark is not used in the baseline simulations. **What we do not claim:** This paper does not establish existence or uniqueness of a full optimal-control MFG equilibrium under intertemporal exponential utility maximization with endogenous filtering on prices. The myopic rule (29) is a bounded-rational approximation; Subsection 4.3 and Subsection 4.4 provide constructive continuous-time comparisons, but a complete verification of existence/uniqueness conditions for those intertemporal MFG couplings would strengthen the theoretical foundation.

3.3 Discrete-Time Myopic Mean-Field Equilibrium (What We Compute)

The numerical experiments in Section VI implement a discrete-time, bounded-rational equilibrium concept consistent with the myopic rule in Section IV. Fix a time grid $t = 0, 1, \dots, T$ (step size $\Delta t = 1$ in the implementation) and consider the N -agent system induced by: (i) fundamental and price updates (Euler discretization of (8)), (ii) private-signal observations (normalized increments $z_{i,t} = (\xi_{i,t+\Delta t} - \xi_{i,t})/\Delta t$) and Kalman belief updates (Appendix A), and (iii) the myopic inventory rule (29) with the per-step stock–flow fixed point (Proposition 4.5).

Definition (myopic mean-field equilibrium, discrete time). Given a pre-trade mean inventory $\bar{x}_{t-}$ and public price p_t , agent i computes a belief state $(\hat{v}_{i,t}, \Sigma_{i,t})$ and chooses a post-trade inventory $x_{i,t}$ according to the myopic rule, which depends on the mean trading rate $\bar{u}_t$. A discrete-time myopic mean-field equilibrium at time t is a pair $(\bar{x}_t, \bar{u}_t)$ such that

$$\bar{x}_t = \frac{1}{N} \sum_{i=1}^N x_{i,t}, \quad \bar{u}_t = \frac{\bar{x}_t - \bar{x}_{t-}}{\Delta t},$$

with individual choices consistent with $\bar{u}_t$. In the implementation the rule is linear in $\bar{u}_t$ and the scalar fixed point is computed in closed form (Proposition 4.5).

3.4 Discrete-Time Well-Posedness (Existence and Uniqueness of Paths)

We formalize the particle implementation as a coupled discrete-time dynamical system on a fixed time grid.

Assumptions. Fix $T \in \mathbb{N}$ and $\Delta t > 0$ and assume: (A1) parameters satisfy $\kappa > 0$, $\phi > 0$, $\gamma > 0$, $\alpha_0 > 0$ and initial conditions $(v_0, p_0, \{\hat{v}_{i,0}, \Sigma_{i,0}\}_{i=1}^N)$ are finite with $\Sigma_{i,0} \geq 0$; (A2) beliefs evolve via the discrete-time Kalman recursion in Appendix A (so $(\hat{v}_{i,t}, \Sigma_{i,t})$ is uniquely defined given $(\hat{v}_{i,t-1}, \Sigma_{i,t-1})$ and the signal); (A3) for each t , the per-step stock–flow map admits a unique fixed point $\bar{x}_t$ (and hence $\bar{u}_t$), e.g., the contraction condition $0 \leq B_t < \Delta t$ in Theorem 4.6 holds; (A4) price and fundamental updates are given by the Euler discretizations of their SDEs driven by specified shock sequences.

Theorem 3.1 (Well-posedness of the discrete-time particle system). *Under (A1)–(A4), for any realization of the common and idiosyncratic shock sequences, the discrete-time system defined by: (i) belief update (Appendix A), (ii) per-step myopic mean-field fixed point (Proposition 4.5), and (iii) Euler updates for (v_t, p_t) , admits a unique adapted solution path*

$$(v_t, p_t, \{\hat{v}_{i,t}, \Sigma_{i,t}\}_{i=1}^N, \bar{x}_t, \bar{u}_t)_{t=0}^T.$$

3.5 Position-impact benchmark: common-noise fixed-point verification (continuous time)

This section is a diagnostic benchmark using the legacy position-impact closure; the baseline simulations use the discrete-time stock–flow equilibrium with uniqueness condition $0 \leq B_t < \Delta t$ (Proposition 4.5, Theorem 4.6).

This subsection records a self-contained fixed-point verification for a continuous-time *position-impact* benchmark (which collapses stock and flow so that $\bar{u}_t \equiv \bar{x}_t$). It is included for comparison and is not used in the baseline simulations, which use the discrete-time stock–flow closure above. The key point is that, in our information simplification (filtering on ξ^i only), the belief state dynamics are *conditionally i.i.d.* given the common filtration, so the mean-field object is the conditional law and the equilibrium closure reduces to an explicit algebraic fixed point.

Common and idiosyncratic noises. Recall the common filtration

$$\mathcal{F}_t^0 := \sigma(W_s, U_s, Z_s : s \leq t),$$

and idiosyncratic signal noises $\{B^i\}_{i \geq 1}$ independent across i and independent of (W, U, Z) .

Finite- N (continuous-time) price-taking system. Consider the N -agent system (continuous time) with a *price-taking* interaction through the average demand $\bar{x}_t^N := \frac{1}{N} \sum_{j=1}^N x_t^{j,N}$:

$$dv_t = \mu_v dt + \sigma_v dW_t, \tag{11}$$

$$d\xi_t^i = v_t dt + \sigma_c dU_t + \sigma_\epsilon dB_t^i, \quad i = 1, \dots, N, \tag{12}$$

$$dp_t^N = \kappa(v_t - p_t^N) dt + \lambda \bar{x}_t^N dt + \phi dZ_t. \tag{13}$$

Agents compute the subjective Kalman–Bucy filter (Appendix A) from (12):

$$\begin{aligned} d\hat{v}_t^i &= \mu_v dt + K_t^{(i)} (d\xi_t^i - \hat{v}_t^i dt), \\ K_t^{(i)} &:= \frac{\Sigma_t^{(i)}}{R_i'}, \quad R_i' := \sigma_c^2 + \sigma_\epsilon^2/k_i. \end{aligned} \tag{14}$$

with Riccati $\dot{\Sigma}_t^{(i)} = \sigma_v^2 - (\Sigma_t^{(i)})^2/R_i'$ and $\Sigma_0^{(i)} \geq 0$. (For a homogeneous population, $k_i \equiv k$ and hence $K_t^{(i)} \equiv K_t$ and $\Sigma_t^{(i)} \equiv \Sigma_t$ are deterministic.)

Given $(p_t^N, \bar{x}_t^N)$, the (benchmark) myopic CARA rule under Gaussian price increments (with $\chi = 0$ and $\text{Var}(dp_t | \mathcal{F}_t^i)/dt = \phi^2$) prescribes the pointwise best response

$$x_t^{i,N} = \frac{\kappa(\hat{v}_t^i - p_t^N) + \lambda \bar{x}_t^N}{\gamma \phi^2}, \quad i = 1, \dots, N. \quad (15)$$

Averaging (15) yields the *per-time* fixed point

$$(\gamma \phi^2 - \lambda) \bar{x}_t^N = \kappa(\hat{m}_t^N - p_t^N), \quad \hat{m}_t^N := \frac{1}{N} \sum_{i=1}^N \hat{v}_t^i, \quad (16)$$

so whenever $\gamma \phi^2 > \lambda$,

$$\bar{x}_t^N = \frac{\kappa}{\gamma \phi^2 - \lambda} (\hat{m}_t^N - p_t^N). \quad (17)$$

Mean-field (conditional-law) system. Define the representative-agent belief state $\hat{v}_t$ solving (14) with an independent copy B (and the same common noises (W, U, Z)), and define the conditional flow

$$\begin{aligned} \mu_t &:= \mathcal{L}(\hat{v}_t, \Sigma_t, k | \mathcal{F}_t^0), \\ m_t &:= \int \hat{v} \mu_t(d\hat{v} d\Sigma dk) = \mathbb{E}[\hat{v}_t | \mathcal{F}_t^0]. \end{aligned}$$

The mean-field price satisfies

$$dp_t = \kappa(v_t - p_t) dt + \lambda \bar{x}_t dt + \phi dZ_t, \quad (18)$$

where $\bar{x}_t$ is required to satisfy the mean-field consistency condition

$$\bar{x}_t = \int \frac{\kappa(\hat{v} - p_t) + \lambda \bar{x}_t}{\gamma \phi^2} \mu_t(d\hat{v} d\Sigma dk). \quad (19)$$

Since only the mean of $\hat{v}$ enters (19), the fixed point reduces to

$$\bar{x}_t = \frac{\kappa}{\gamma \phi^2 - \lambda} (m_t - p_t), \quad \text{assuming } \gamma \phi^2 > \lambda. \quad (20)$$

Theorem 3.2 (Common-noise fixed point). *Assume $\gamma \phi^2 > \lambda$ and $\sup_{i \leq N} \mathbb{E}[|\hat{v}_0^i|^2] < \infty$ with $\sup_{i \leq N} \Sigma_0^{(i)} < \infty$. For simplicity, assume initial belief states are i.i.d. across i conditional on $\mathcal{F}_0^0$. Then:*

(i) (Well-posedness.) *The N -agent system (11)–(13) with controls given by the myopic fixed point (15) admits a unique strong solution, and the limiting mean-field (conditional-law) system consisting of (11), a representative signal equation of the form (12) (with an independent idiosyncratic Brownian motion), the representative filter (14), and the price equation (18) with $\bar{x}$ given by (20) admits a unique strong solution.*

(ii) (Fixed-point verification.) *The process $\bar{x}$ defined by (20) is the unique solution to the mean-field consistency equation (19) (i.e., a unique myopic common-noise MFG fixed point).*

(iii) (Conditional propagation of chaos / LLN.) *Let the empirical measure of belief states be $\mu_t^N := \frac{1}{N} \sum_{i=1}^N \delta_{(\hat{v}_t^i, \Sigma_t^{(i)}, k_i)}$. Then for each fixed t ,*

$$\mu_t^N \xrightarrow[N \rightarrow \infty]{} \mu_t \quad \text{in probability, and} \quad \hat{m}_t^N \rightarrow m_t \text{ in } L^2.$$

Consequently, $\bar{x}_t^N \rightarrow \bar{x}_t$ in L^2 and $p_t^N \rightarrow p_t$ in L^2 uniformly on $[0, T]$.

Proof. Step 1 (Well-posedness). For each i , the filter SDE (14) is linear with adapted coefficients and admits a unique strong solution with finite second moments; the Riccati equation is a scalar ODE with global solution (Appendix A). Given $\bar{x}^N$ and p^N , the control (15) is affine in $(\hat{v}^i, p^N, \bar{x}^N)$. The averaging identity (16) yields the explicit closure (17) when $\gamma\phi^2 > \lambda$. Substituting (17) into (13) gives a linear SDE for p^N driven by (W, U, Z) and the averaged belief $\hat{m}^N$, hence a unique strong solution. Likewise, substituting (20) into (18) gives a linear SDE for p driven by (W, U, Z) and m .

Step 2 (Fixed point). Equation (19) is affine in $\bar{x}_t$:

$$\bar{x}_t = \frac{\kappa(m_t - p_t) + \lambda\bar{x}_t}{\gamma\phi^2}.$$

Rearranging gives $(\gamma\phi^2 - \lambda)\bar{x}_t = \kappa(m_t - p_t)$ and thus (20). Uniqueness follows since $\gamma\phi^2 - \lambda > 0$.

Step 3 (Conditional LLN / propagation of chaos). Conditional on $\mathcal{F}_t^0$, the idiosyncratic noises $\{B^i\}$ are independent and the filter equations (14) are driven by independent Brownian motions with the *same* common coefficients and common path (W, U, Z) . Hence the collection $\{(\hat{v}_t^i, \Sigma_t^{(i)}, k_i)\}_{i=1}^N$ is i.i.d. conditional on $\mathcal{F}_t^0$ (with law μ_t), and the conditional law of large numbers implies $\mu_t^N \Rightarrow \mu_t$ and

$$\mathbb{E}[|\hat{m}_t^N - m_t|^2 \mid \mathcal{F}_t^0] = \frac{1}{N} \text{Var}(\hat{v}_t \mid \mathcal{F}_t^0),$$

so $\hat{m}_t^N \rightarrow m_t$ in L^2 . Moreover, by Step 1 we have $\sup_{t \leq T} \mathbb{E}[\hat{v}_t^2] < \infty$, hence $\sup_{t \leq T} \mathbb{E}[\text{Var}(\hat{v}_t \mid \mathcal{F}_t^0)] \leq \sup_{t \leq T} \mathbb{E}[\hat{v}_t^2] < \infty$ and therefore

$$\sup_{t \leq T} \mathbb{E}|\hat{m}_t^N - m_t|^2 \leq \frac{C_T}{N}, \quad \int_0^T \mathbb{E}|\hat{m}_s^N - m_s|^2 ds \leq \frac{TC_T}{N} \xrightarrow{N \rightarrow \infty} 0.$$

Step 4 (Price convergence). Let $C := \kappa/(\gamma\phi^2 - \lambda)$ so that (17)–(20) read $\bar{x}_t^N = C(\hat{m}_t^N - p_t^N)$ and $\bar{x}_t = C(m_t - p_t)$. Substituting these into (13) and (18) gives the closed price equations

$$dp_t^N = \kappa(v_t - p_t^N) dt + \lambda C(\hat{m}_t^N - p_t^N) dt + \phi dZ_t, \quad dp_t = \kappa(v_t - p_t) dt + \lambda C(m_t - p_t) dt + \phi dZ_t.$$

Subtracting yields a linear ODE for $\Delta p_t := p_t^N - p_t$ driven by $\Delta m_t := \hat{m}_t^N - m_t$:

$$d\Delta p_t = -(\kappa + \lambda C)\Delta p_t dt + \lambda C \Delta m_t dt.$$

By variation of constants, for $t \in [0, T]$,

$$\Delta p_t = e^{-(\kappa + \lambda C)t} \Delta p_0 + \lambda C \int_0^t e^{-(\kappa + \lambda C)(t-s)} \Delta m_s ds.$$

Assuming matching initial conditions ($\Delta p_0 = 0$), Jensen and Cauchy–Schwarz imply

$$\sup_{t \leq T} \mathbb{E}|p_t^N - p_t|^2 \leq (\lambda C)^2 T \int_0^T \mathbb{E}|\hat{m}_s^N - m_s|^2 ds \xrightarrow{N \rightarrow \infty} 0,$$

and the L^2 uniform convergence follows.

Step 5 (Mean-field action convergence). Using again $\bar{x}_t^N = C(\hat{m}_t^N - p_t^N)$ and $\bar{x}_t = C(m_t - p_t)$,

$$\bar{x}_t^N - \bar{x}_t = C[(\hat{m}_t^N - m_t) - (p_t^N - p_t)].$$

Thus, by Step 3 and Step 4,

$$\sup_{t \leq T} \mathbb{E}|\bar{x}_t^N - \bar{x}_t|^2 \leq 2C^2 \left(\sup_{t \leq T} \mathbb{E}|\hat{m}_t^N - m_t|^2 + \sup_{t \leq T} \mathbb{E}|p_t^N - p_t|^2 \right) \xrightarrow{N \rightarrow \infty} 0.$$

□

Remark 3.3 (Stochastic Fokker–Planck equation for μ_t). For any test function $\varphi \in C_b^2(\mathbb{R} \times \mathbb{R}_+ \times \mathbb{R}_+)$, the conditional law $\mu_t = \mathcal{L}(\hat{v}_t, \Sigma_t, k \mid \mathcal{F}_t^0)$ satisfies the common-noise SPDE (in weak form)

$$\begin{aligned} d\langle \mu_t, \varphi \rangle &= \left\langle \mu_t, (\mu_v + K_t(k)(v_t - \hat{v})) \partial_{\hat{v}} \varphi \right. \\ &\quad \left. + \frac{1}{2} K_t(k)^2 \sigma_\epsilon^2 \partial_{\hat{v}\hat{v}}^2 \varphi \right\rangle dt \\ &\quad + \left\langle \mu_t, K_t(k) \sigma_c \partial_{\hat{v}} \varphi \right\rangle dU_t. \end{aligned}$$

where $\langle \mu, \varphi \rangle := \int \varphi d\mu$ and $K_t(k) = \Sigma_t(k)/(\sigma_c^2 + \sigma_\epsilon^2/k)$. In the homogeneous case, μ_t is Gaussian in $\hat{v}$ conditional on $\mathcal{F}_t^0$ (with random mean m_t).

3.6 Mean-Field Limit

We state a standard common-noise McKean–Vlasov/MFG limit as a black-box framework; our contribution is the constructive solution in the LQG specialization, given in Section 3.7.

Setting. Consider a common-noise McKean–Vlasov system on a finite horizon $[0, T]$. Let $\mathcal{F}_t^0 := \sigma(W_s^0 : s \leq t)$ be the filtration generated by a common Brownian motion W^0 , and for each $N \geq 1$ let $\{B^i\}_{i=1}^N$ be mutually independent and independent of W^0 . The N -agent state processes $\{X_t^{i,N}\}_{i=1}^N$ take values in $\mathbb{R}^d$ and satisfy a coupled system whose coefficients may depend on the empirical measure $\bar{\mu}_t^N := \frac{1}{N} \sum_{j=1}^N \delta_{X_t^{j,N}}$ and on the common noise. The mean-field limit is described by a McKean–Vlasov SDE in which the conditional law $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^0)$ of a representative state X_t enters the drift and diffusion; see Carmona and Delarue (2018a,b); Lacker (2018).

Assumptions. We impose the following standard conditions (see, e.g., Fournier and Guillin (2015); Lacker (2018)).

- (H1) (**Lipschitz in state and measure.**) The drift $b(t, x, \mu)$ and diffusion $\sigma(t, x, \mu)$ are Lipschitz in (x, μ) uniformly in $t \in [0, T]$, with μ in the space of probability measures on $\mathbb{R}^d$ metrized by the Wasserstein distance W_2 .
- (H2) (**Linear growth.**) $|b(t, x, \mu)| + \|\sigma(t, x, \mu)\| \leq C(1 + |x| + W_2(\mu, \delta_0))$ for some $C \geq 0$.
- (H3) (**Common noise integrability.**) Common-noise coefficients are progressively measurable and square-integrable on $[0, T]$ a.s.
- (H4) (**Initial condition.**) The initial states $\{X_0^{i,N}\}_{i=1}^N$ are i.i.d. with law μ_0 and satisfy $\mathbb{E}[|X_0^{i,N}|^2] < \infty$; μ_0 has finite second moment.
- (H5) (**Non-degeneracy / bounded coefficients.**) The diffusion matrix $\sigma\sigma^\top$ is uniformly elliptic (or at least non-degenerate) and all coefficients are bounded when restricted to compacts in $\mathbb{R}^d \times \mathcal{P}_2(\mathbb{R}^d)$.

Under (H1)–(H5), the limiting McKean–Vlasov equation admits a unique strong solution and the N -particle system is well-posed; see Carmona and Delarue (2018a,b).

Theorem 3.4 (Convergence of empirical measures). *Assume (H1)–(H5). Let $\{X^{i,N}\}_{i=1}^N$ solve the N -particle system and let X solve the limiting McKean–Vlasov equation with the same common noise W^0 and i.i.d. copies $\{B^i\}$ of the idiosyncratic driver. Then for each $t \in [0, T]$,*

$$\bar{\mu}_t^N \xrightarrow[N \rightarrow \infty]{W_2} \mu_t \quad \text{in probability,}$$

where $\mu_t = \mathcal{L}(X_t \mid \mathcal{F}_t^0)$ is the conditional law of the limit process. In particular, propagation of chaos holds: any fixed finite marginals converge in law to the product of the limit conditional law.

Proof sketch (black-box references). Existence and uniqueness for the limit equation follow from (H1)–(H2) and fixed-point arguments in the space of flows of measures; see Carmona and Delarue (2018a). Convergence of $\bar{\mu}_t^N$ to μ_t in W_2 is a standard consequence of (H1)–(H5) and the conditional law of large numbers for the empirical measure; see Fournier and Guillin (2015) for quantitative rates and Lacker (2018) for the common-noise formulation. Propagation of chaos is then implied by the same estimates; see Carmona and Delarue (2018b). $\square$

Proposition 3.5 (Stability of mean-field equilibrium). *Under (H1)–(H5), the mean-field equilibrium flow $\{\mu_t\}_{t \in [0, T]}$ is unique in the class of admissible flows with finite second moments. Moreover, $\sup_{t \in [0, T]} \mathbb{E}[W_2(\bar{\mu}_t^N, \mu_t)^2] \leq C/N$ for some constant C depending on T and the data.*

Proof sketch (black-box references). Uniqueness follows from the Lipschitz structure (H1) and Gronwall-type estimates in Wasserstein space; see Carmona and Delarue (2018a). The $O(1/N)$ bound for $\mathbb{E}[W_2(\bar{\mu}_t^N, \mu_t)^2]$ is given in Fournier and Guillin (2015) and Huang et al. (2006) for the classical and MFG settings. $\square$

Remark 3.6 (LQG specialization). In our LQG setting (see Section IV), the coefficients are affine in state and measure, so (H1)–(H2) hold and the limit is explicit. The preceding black-box statements justify using the mean-field system as the large- N approximation for the particle implementation in Section VI.

3.7 Constructive LQG Solution via Picard Iteration

Benchmark (position impact; continuous time). This subsection is a continuous-time constructive solver for a *position-impact* formulation (it collapses stock and flow so that $\bar{u} \equiv \bar{x}$). It is included as a diagnostic/benchmark and is not used in the baseline discrete-time stock–flow simulations in Section VI.

Endogenous vs exogenous. Mean demand $\bar{x}$, price p , and (from the equilibrium map) mispricing $\bar{y} = m - p$ and the coefficient B_t (driven by $\bar{x}$) are *endogenous*; in equilibrium, beliefs and the filter are consistent with the chosen information. Exogenous parameters include $\kappa, \lambda, \phi, \gamma$, and other primitives (e.g. $\sigma_v, \sigma_c, \sigma_\epsilon, k$). In this formulation, belief updating conditions only on the private signal ξ and not on the price observation channel; hence the mean belief m_t is not fed back by $\bar{x}$. Price is endogenous as an outcome of $\bar{x}$ but is not used as an observation in the filter; incorporating it would require joint filtering and a richer fixed-point formulation (see Section IV).

Reduced state and dynamics. Define the (belief–price) mispricing state

$$y_t := \hat{v}_t - p_t. \quad (21)$$

We work in the homogeneous representative-agent setting with overconfidence parameter k , so $\hat{v}_t$ is the subjective Kalman estimate of v_t , Σ_t is the perceived posterior variance, and $R' := \sigma_c^2 + \sigma_\epsilon^2/k$ is the perceived measurement-noise variance. Wealth and state dynamics in innovation form are

$$dW_t = x_t (\kappa y_t + \lambda \bar{x}_t) dt + x_t \phi dI_t^p, \quad (22)$$

$$dy_t = (\mu_v - \kappa y_t - \lambda \bar{x}_t) dt + \beta_t dI_t^\xi - \phi dI_t^p, \quad (23)$$

where (under the agent's subjective filtering model in Appendix A) I^p and I^ξ are independent Brownian motions (price and signal innovations). The (subjective) unhedgeable loading is

$$\beta_t = K_t \sqrt{\sigma_c^2 + \sigma_\epsilon^2}, \quad K_t = \frac{\Sigma_t}{R'}, \quad R' = \sigma_c^2 + \frac{\sigma_\epsilon^2}{k}, \quad (24)$$

and the perceived posterior variance Σ_t solves the Riccati equation

$$\dot{\Sigma}_t = \sigma_v^2 - \frac{\Sigma_t^2}{R'}, \quad \Sigma_0 \geq 0. \quad (25)$$

Fixed-point map and contraction. Theorem 3.4 provides the framework: existence, uniqueness, and convergence of the mean-field limit are from the black-box theory (assumptions (H1)–(H5)); the reduced dynamics (22)–(23), Riccati (25), and mean-field update (26) below are derived in our LQG setting. Given a candidate mean-field demand process $\bar{x}(\cdot)$, one can (i) solve the backward Riccati for A_t (independent of $\bar{x}$), (ii) solve the backward linear ODE for B_t driven by $\bar{x}_t$, (iii) simulate forward the coupled (v_t, m_t, p_t) and hence $\bar{y}_t = m_t - p_t$, and (iv) update $\bar{x}$ via the mean-field consistency relation. This defines a fixed-point map Φ on trajectories $\bar{x}(\cdot)$.

Fixed-point object. The fixed-point object is the map Φ on the space $\mathcal{X} := L_{\mathcal{F}^0}^2(\Omega \times [0, T])$ of candidate mean-field demand trajectories (adapted to the common filtration $\mathcal{F}^0$), equipped with the norm

$$\|\bar{x}\|_2 := \left(\mathbb{E} \int_0^T |\bar{x}_t|^2 dt \right)^{1/2}.$$

Equilibrium corresponds to the unique fixed point $\bar{x}^* = \Phi(\bar{x}^*)$ when the contraction condition $q < 1$ below holds; see Proposition 3.7.

Proposition 3.7 (Constructive equilibrium under weak coupling / small horizon). *Define the fixed-point map Φ on $\bar{x}(\cdot)$ as follows. Given $\bar{x}(\cdot)$:*

- (i) *Solve backward the Riccati equation for A_t with terminal condition $A_T = 0$ (as in the A-ODE of the exact CARA/LQG formulation).*
- (ii) *Solve backward the linear ODE for B_t driven by $\bar{x}_t$, with terminal condition $B_T = 0$.*
- (iii) *Solve forward the dynamics for (v_t, m_t, p_t) given $\bar{x}_t$, and set $\bar{y}_t = m_t - p_t$.*
- (iv) *Update via mean-field consistency:*

$$\bar{x}_t = \frac{(\kappa + \phi^2 A_t) \bar{y}_t + \phi^2 B_t}{\gamma \phi^2 - \lambda}. \quad (26)$$

Define

$$q := \frac{\lambda T}{\gamma \phi^2 - \lambda} (2\kappa + \kappa^2 T). \quad (27)$$

If $\gamma \phi^2 > \lambda$ and $q < 1$, then Φ is a contraction on the space $\mathcal{X} := L_{\mathcal{F}^0}^2(\Omega \times [0, T])$ and the fixed point exists and is unique; Picard iteration converges.

Proof. We make the contraction argument fully explicit.

Functional-analytic setup. Fix a filtered probability space $(\Omega, \mathcal{F}, \mathcal{F}^0, \mathbb{P})$ supporting the common shocks (e.g. (W, U, Z) as in (32)–(34)) and the cross-sectional private noises. Let

$\mathcal{F}^0 = (\mathcal{F}_t^0)_{t \in [0, T]}$ be the (common) market filtration containing the common shocks. Define the Banach space

$$\mathcal{X} := L^2_{\mathcal{F}^0}(\Omega \times [0, T]), \quad \|\bar{x}\|_2 := \left(\mathbb{E} \int_0^T |\bar{x}_t|^2 dt \right)^{1/2}.$$

(If one prefers to restrict to deterministic mean-field trajectories $\bar{x} : [0, T] \rightarrow \mathbb{R}$, the same proof applies verbatim with $\|\bar{x}\|_{L^2([0, T])} := (\int_0^T |\bar{x}_t|^2 dt)^{1/2}$.)

Assume $\gamma\phi^2 > \lambda$, so the denominator in (26) is strictly positive. Also assume β is $\mathcal{F}^0$ -progressively measurable and bounded on $[0, T]$ (this holds under the constant-gain closure (31), or any bounded-gain filtering specification used in Sections III–IV).

Step 1: Riccati well-posedness and a uniform bound for A . In the homogeneous case, A_t solves (44) with terminal condition $A_T = 0$:

$$-\dot{A}_t = \frac{\kappa^2}{\phi^2} - \beta_t^2 A_t^2, \quad A_T = 0.$$

Set $\tilde{A}_s := A_{T-s}$ for $s \in [0, T]$. Then

$$\dot{\tilde{A}}_s = \frac{\kappa^2}{\phi^2} - \beta_{T-s}^2 \tilde{A}_s^2, \quad \tilde{A}_0 = 0.$$

Nonnegativity. The vector field $f(s, a) := \kappa^2/\phi^2 - \beta_{T-s}^2 a^2$ is continuous in s and locally Lipschitz in a , hence solutions are unique. Moreover $f(s, 0) = \kappa^2/\phi^2 > 0$ for all s , so starting from $\tilde{A}_0 = 0$ the solution becomes strictly positive immediately. If it ever became negative, let $\tau := \inf\{s > 0 : \tilde{A}_s = 0\}$ be the first return time to 0 from above. Then $\dot{\tilde{A}}_\tau = f(\tau, 0) > 0$, so the trajectory points into the positive half-line at τ , contradicting the definition of τ . Thus $\tilde{A}_s \geq 0$ for all $s \in [0, T]$, hence $A_t \geq 0$ for all t . *Coarse bound.* Because the quadratic term is nonnegative,

$$\dot{\tilde{A}}_s \leq \frac{\kappa^2}{\phi^2} \implies 0 \leq \tilde{A}_s \leq \frac{\kappa^2}{\phi^2} s,$$

hence

$$0 \leq A_t \leq \frac{\kappa^2}{\phi^2}(T - t) \leq \frac{\kappa^2}{\phi^2}T \quad \text{for all } t \in [0, T].$$

In particular, $A \in L^\infty([0, T])$ and $\|A\|_\infty \leq \kappa^2 T / \phi^2$.

Step 2: Precise assumption on mean belief and the $\bar{y}$ -Lipschitz bound. Recall $\bar{y}_t := m_t - p_t$ in Proposition III.7.

In the myopic formulation of Section IV.B, belief updating (Appendix A) is conditioned only on the private signal $\xi_{i,t}$ (and not on the endogenous price observation channel). Under this modeling choice, the mean belief process $m_t = \int_0^1 \hat{v}_{i,t} di$ is independent of the mean-field demand input $\bar{x}$; i.e. for two inputs $\bar{x}, \bar{x}' \in \mathcal{X}$ we have $m \equiv m'$.

(If one prefers not to hard-code this, it suffices to assume instead that $\bar{x} \mapsto m$ is Lipschitz in $\|\cdot\|_2$ with constant $C_m T$; the bound below then holds with λ replaced by $\lambda + C_m$.)

Now fix $\bar{x}, \bar{x}' \in \mathcal{X}$. Couple the two forward systems on the same probability space with the same common shocks and identical initial conditions (v_0, p_0) . Subtracting the price equations (34) yields, pathwise,

$$d(p_t - p'_t) = -\kappa(p_t - p'_t) dt + \lambda(\bar{x}_t - \bar{x}'_t) dt, \quad (p_0 - p'_0) = 0,$$

since the common noise term ϕdZ_t cancels under this coupling. Solving gives

$$|p_t - p'_t| \leq \lambda \int_0^t e^{-\kappa(t-s)} |\bar{x}_s - \bar{x}'_s| ds \leq \lambda \int_0^t |\bar{x}_s - \bar{x}'_s| ds \leq \lambda \sqrt{t} \left(\int_0^t |\bar{x}_s - \bar{x}'_s|^2 ds \right)^{1/2}.$$

Squaring and integrating in t , then using Fubini, yields

$$\int_0^T |p_t - p'_t|^2 dt \leq \lambda^2 \int_0^T t \int_0^t |\bar{x}_s - \bar{x}'_s|^2 ds dt \leq \lambda^2 T^2 \int_0^T |\bar{x}_s - \bar{x}'_s|^2 ds.$$

Taking expectations gives $\|p - p'\|_2 \leq \lambda T \|\bar{x} - \bar{x}'\|_2$. Because $m \equiv m'$, we have $\bar{y} - \bar{y}' = -(p - p')$, hence

$$\|\bar{y} - \bar{y}'\|_2 \leq \lambda T \|\bar{x} - \bar{x}'\|_2.$$

Step 3: B -Lipschitz bound in $\bar{x}$. Given A (which does not depend on $\bar{x}$) and β (independent of $\bar{x}$ under the same information closure as above), B_t solves (45) with terminal condition $B_T = 0$:

$$-\dot{B}_t = \frac{\kappa\lambda}{\phi^2} \bar{x}_t + \mu_v A_t - \beta_t^2 A_t B_t, \quad B_T = 0.$$

(If β is allowed to depend on $\bar{x}$, an additional Lipschitz term appears; the same contraction logic goes through with a larger constant.)

Let B, B' be the solutions corresponding to $\bar{x}, \bar{x}'$ and set $\Delta B := B - B'$. Subtracting the ODEs gives

$$-\dot{\Delta B}_t = \frac{\kappa\lambda}{\phi^2} (\bar{x}_t - \bar{x}'_t) - \beta_t^2 A_t \Delta B_t, \quad \Delta B_T = 0.$$

By integrating factor,

$$\Delta B_t = \frac{\kappa\lambda}{\phi^2} \int_t^T \exp\left(-\int_t^s \beta_u^2 A_u du\right) (\bar{x}_s - \bar{x}'_s) ds.$$

Since $A_u \geq 0$ and $\beta_u^2 \geq 0$, the exponential factor is at most 1, yielding

$$|B_t - B'_t| \leq \frac{\kappa\lambda}{\phi^2} \int_t^T |\bar{x}_s - \bar{x}'_s| ds \leq \frac{\kappa\lambda}{\phi^2} \sqrt{T-t} \left(\int_t^T |\bar{x}_s - \bar{x}'_s|^2 ds \right)^{1/2}.$$

Squaring and integrating in t , then using Fubini, yields $\|B - B'\|_2 \leq (\kappa\lambda/\phi^2) T \|\bar{x} - \bar{x}'\|_2$.

Step 4: Contraction estimate for Φ and a fully checkable sufficient condition. Recall (26):

$$\Phi(\bar{x})_t = \frac{(\kappa + \phi^2 A_t) \bar{y}_t + \phi^2 B_t}{\gamma\phi^2 - \lambda}.$$

Using $\gamma\phi^2 - \lambda > 0$, boundedness of A , and the triangle inequality,

$$\|\Phi(\bar{x}) - \Phi(\bar{x}')\|_2 \leq \frac{(\kappa + \phi^2 \|A\|_\infty) \|\bar{y} - \bar{y}'\|_2 + \phi^2 \|B - B'\|_2}{\gamma\phi^2 - \lambda}.$$

Substituting the bounds from Steps 2–3 gives

$$\|\Phi(\bar{x}) - \Phi(\bar{x}')\|_2 \leq \frac{\lambda T}{\gamma\phi^2 - \lambda} \left(2\kappa + \phi^2 \|A\|_\infty \right) \|\bar{x} - \bar{x}'\|_2.$$

Finally, using $\|A\|_\infty \leq \kappa^2 T / \phi^2$ from Step 1 yields the parameter-only bound

$$\|\Phi(\bar{x}) - \Phi(\bar{x}')\|_2 \leq q \|\bar{x} - \bar{x}'\|_2, \quad q := \frac{\lambda T}{\gamma\phi^2 - \lambda} \left(2\kappa + \kappa^2 T \right),$$

Algorithm 1 Picard iteration for LQG mean-field equilibrium

- 1: **Input:** Horizon T , parameters $(\kappa, \lambda, \gamma, \phi, \mu_v, \sigma_v, \sigma_c, \sigma_\epsilon, k)$, initial conditions $(v_0, p_0, \hat{v}_0, \Sigma_0)$, tolerance ϵ , max iterations $N_{\max}$.
 - 2: **Initialize:** Set $\bar{x}_t^0 \leftarrow 0$ (or myopic seed) for $t \in [0, T]$. Set $n \leftarrow 0$.
 - 3: **repeat**
 - 4: Compute Σ_t, K_t, β_t on $[0, T]$ from (25) and (24).
 - 5: Solve backward for A_t (Riccati, $A_T = 0$).
 - 6: Solve backward for B_t given $\bar{x}_t^n$ (linear ODE, $B_T = 0$).
 - 7: Forward-simulate (v_t, m_t, p_t) and $\bar{y}_t^n = m_t - p_t$ using $\bar{x}_t^n$ in the dynamics.
 - 8: Update $\bar{x}_t^{n+1} \leftarrow \frac{(\kappa + \phi^2 A_t) \bar{y}_t^n + \phi^2 B_t}{\gamma \phi^2 - \lambda}$ for $t \in [0, T]$.
 - 9: Compute residual $r_n \leftarrow \|\bar{x}^{n+1} - \bar{x}^n\|$ (e.g., L^2 or sup over grid).
 - 10: $n \leftarrow n + 1$.
 - 11: **until** $r_{n-1} < \epsilon$ or $n \geq N_{\max}$
 - 12: **Output:** Equilibrium mean-field demand $\bar{x}^* \approx \bar{x}^n$ and associated $(A, B, \bar{y})$.
-

which is fully checkable from primitives.

Step 5: Banach fixed point. If $q < 1$, then Φ is a contraction on $(\mathcal{X}, \|\cdot\|_2)$, hence admits a unique fixed point $\bar{x}^* \in \mathcal{X}$. Moreover, the Picard iteration $\bar{x}^{n+1} = \Phi(\bar{x}^n)$ converges geometrically to $\bar{x}^*$ in $\|\cdot\|_2$.

Remark (relation to the simpler myopic scaling). In the myopic closure where hedging terms are suppressed ($A \equiv 0, B \equiv 0$), the Lipschitz estimate tightens to a pure $O(T)$ scaling proportional to $\kappa\lambda/(\gamma\phi^2 - \lambda)$, matching the heuristic small-horizon/weak-coupling condition discussed around (27). $\square$

Picard iteration. A conservative Lipschitz bound is q in (27); for $q < 1$, Banach yields uniqueness. Algorithm 1 summarizes the procedure. We initialize a candidate $\bar{x}^0(\cdot)$ (e.g., identically zero or the myopic closure $\bar{x}_t^0 = \kappa\bar{y}_t^0/(\gamma\phi^2 - \lambda)$ from a forward run with $\bar{x} \equiv 0$), then iterate: compute Σ, K, β from (25)–(24), solve A backward, solve B backward using the current $\bar{x}^n$, forward-simulate (v, m, p) and $\bar{y}^n$, update $\bar{x}^{n+1}$ via (26), and stop when the residual (e.g., $\|\bar{x}^{n+1} - \bar{x}^n\|$ in L^2 or sup-norm over the time grid) is below a tolerance.

Implementation and reproducibility. The Picard iteration in Algorithm 1 is implemented in a MATLAB script; the path and run instructions will be specified in the replication package. All reported equilibrium paths and statistics use this routine with a fixed tolerance and time discretization.

3.8 Reproducibility and numerical verification

The equilibrium is computed via Algorithm 1 implemented in `code/matlab/mfg_lqg_common_noise_demo.m`. Diagnostics include the fixed-point residual along Picard iterates, a time-step refinement test under decreasing dt , and an N -convergence test comparing particle moments to the mean-field limit.²

Remark (discrete-time certification). At the discrete-time level, uniqueness of the numerical fixed point can be certified via a Banach fixed-point argument on trajectories endowed with a sup-norm metric. We provide a machine-checkable backbone for this contraction argument in

²Outputs and plots can be regenerated by running the script.

lean/MFGContraction.lean; it does not attempt to formalize the full stochastic propagation-of-chaos theorem.

4 Main Results

4.1 Static Benchmark: Price Sensitivity and Overconfidence

The closed-form static price (5) yields explicit comparative statics with respect to overconfidence.

Proposition 4.1 (Static price response increases with overconfidence). *Let $\tau_{v,0} := 1/\sigma_{v,0}^2$ and $\tau'_{\epsilon,0} := k/\sigma_{\epsilon,0}^2$. In the static benchmark, the coefficient on v in (5) is*

$$a(k) = \frac{\lambda\tau'_{\epsilon,0}}{\gamma + \lambda(\tau_{v,0} + \tau'_{\epsilon,0})}, \quad (28)$$

which is increasing in k and satisfies $0 < a(k) < 1$.

4.2 Myopic Trading Rule and Mean-Field Consistency

We next derive a checkable *myopic* inventory rule for the discrete-time timing in Subsubsection 2.3.3. Fix a decision interval Δt and treat the mean trading rate $\bar{u}_t$ and the public price p_t as given. For tractability, belief updating (Appendix A) conditions only on the private signal ξ_i ; incorporating endogenous prices as an additional observation channel would require joint filtering and a richer fixed-point formulation. Conditional on the investor's information $\mathcal{F}_t^i$, the one-step price increment $\Delta p_{t+\Delta t}$ is Gaussian with conditional mean $(\kappa(\hat{v}_{i,t} - p_t) + \lambda\bar{u}_t)\Delta t$ and conditional variance $\sigma_{p,i,t}^2\Delta t$.

Proposition 4.2 (Myopic CARAinventory choice with holding cost and flow impact). *Fix a decision interval Δt . Suppose conditional on $\mathcal{F}_t^i$ the price increment is Gaussian with $\mathbb{E}_i[\Delta p_{t+\Delta t}] = (\kappa(\hat{v}_{i,t} - p_t) + \lambda\bar{u}_t)\Delta t$ and $\text{Var}_i(\Delta p_{t+\Delta t}) = \sigma_{p,i,t}^2\Delta t$. With one-step wealth $\Delta W_{i,t} = x_{i,t}\Delta p_{t+\Delta t} - (\chi/2)x_{i,t}^2\Delta t$ and CARAutility, the unique maximizer is*

$$x_{i,t}^* = \frac{\kappa(\hat{v}_{i,t} - p_t) + \lambda\bar{u}_t}{\chi + \gamma\sigma_{p,i,t}^2}. \quad (29)$$

Define the corresponding *mispricing gain* (shares per unit price) as the coefficient on $(\hat{v}_{i,t} - p_t)$ in (29). Under the Euler increment (8) over Δt and conditioning on $\mathcal{F}_t^i$, $\sigma_{p,i,t}^2 = \phi^2 + \kappa^2\Sigma_{i,t}\Delta t$, so

$$\Theta_t(\Sigma_{i,t}) := \frac{\kappa}{\chi + \gamma\sigma_{p,i,t}^2} = \frac{\kappa}{\chi + \gamma(\phi^2 + \kappa^2\Sigma_{i,t}\Delta t)}. \quad (30)$$

Moreover,

$$\partial_\Sigma\Theta_t(\Sigma) = -\frac{\gamma\kappa^3\Delta t}{(\chi + \gamma(\phi^2 + \kappa^2\Sigma\Delta t))^2} \leq 0,$$

so Θ_t is weakly decreasing in Σ .

Position-impact benchmark (reduced form; not used in baseline simulations). If one collapses stock and flow (so $\bar{u}_t \equiv \bar{x}_t$) and sets $\chi = 0$ and $\sigma_{p,i,t}^2 \equiv \phi^2$, then averaging (29) yields the familiar homogeneous closure (assuming $\gamma\phi^2 > \lambda$):

$$\bar{x}_t = \frac{\kappa}{\gamma\phi^2 - \lambda}(m_t - p_t), \quad (31)$$

where $m_t := \int_0^1 \hat{v}_{i,t} di$ is the mean belief. This reduced form is retained only as a transparent diagnostic benchmark; the baseline simulations solve the discrete-time stock–flow fixed point with inventory carryover and uniqueness condition $0 \leq B_t < \Delta t$ (Proposition 4.5 and Theorem 4.6).

4.3 Exact continuous-time CARA/LQG control with filtering and mean-field coupling

This subsection replaces the infinitesimal-horizon (myopic) CARA rule with the finite-horizon continuous-time exponential-utility optimum under partial information. We keep the same fundamental/signal/price dynamics (32)–(34) and the same Kalman–Bucy filtering structure for subjective beliefs.

Dynamics and information. Let the fundamental and price evolve as

$$dv_t = \mu_v dt + \sigma_v dW_t, \quad (32)$$

$$d\xi_{i,t} = v_t dt + \sigma_c dU_t + \sigma_\epsilon dB_{i,t}, \quad (33)$$

$$dp_t = \kappa(v_t - p_t) dt + \lambda \bar{u}_t dt + \phi dZ_t, \quad (34)$$

where $\bar{u}_t := \int_0^1 u_{j,t} dj$ is the mean-field trading rate (taken as exogenous by an infinitesimal agent). Agent i observes $(p_s, \xi_{i,s})_{s \leq t}$ and forms a subjective filtered estimate $\hat{v}_{i,t}$.

Subjective Kalman filter (with overconfidence). As in Appendix A, agent i runs a (possibly misspecified) Kalman–Bucy filter

$$\begin{aligned} d\hat{v}_{i,t} &= \mu_v dt + K_t^{(i)} (d\xi_{i,t} - \hat{v}_{i,t} dt), \\ K_t^{(i)} &:= \frac{\Sigma_t^{(i)}}{R_i'}, \quad R_i' := \sigma_c^2 + \frac{\sigma_\epsilon^2}{k_i}, \end{aligned} \quad (35)$$

with Riccati equation

$$\dot{\Sigma}_t^{(i)} = \sigma_v^2 - \frac{(\Sigma_t^{(i)})^2}{R_i'}, \quad \Sigma_0^{(i)} \geq 0. \quad (36)$$

(When convenient, one may use the constant-gain approximation $K_t^{(i)} \approx K_i^\star = \sigma_v / \sqrt{R_i'}$.)

Innovation form and reduced state. Define the (belief-based) mispricing state

$$y_{i,t} := \hat{v}_{i,t} - p_t. \quad (37)$$

Under the agent’s filtration, write the observed price in innovations form

$$dp_t = [\kappa y_{i,t} + \lambda \bar{u}_t] dt + \phi dI_{i,t}^p, \quad (38)$$

where I_i^p is a Brownian motion (price innovation). Likewise, the (normalized) signal innovation process is defined by

$$dI_{i,t}^\xi := \frac{d\xi_{i,t} - \hat{v}_{i,t} dt}{\sqrt{\sigma_c^2 + \sigma_\epsilon^2}}.$$

Under the agent filtration $\mathcal{F}_t^i$, I_i^ξ is a Brownian motion and is independent of the price innovation I_i^p .

Consequently, the controlled wealth and the state satisfy

$$dW_{i,t} = x_{i,t} dp_t = x_{i,t} [\kappa y_{i,t} + \lambda \bar{u}_t] dt + x_{i,t} \phi dI_{i,t}^p, \quad (39)$$

$$dy_{i,t} = [\mu_v - \kappa y_{i,t} - \lambda \bar{u}_t] dt + \beta_t^{(i)} dI_{i,t}^\xi - \phi dI_{i,t}^p, \quad (40)$$

where I_i^ξ is an (independent) Brownian motion and the (subjective) unhedgeable belief-noise loading is

$$\beta_t^{(i)} := K_t^{(i)} \sqrt{\sigma_c^2 + \sigma_\epsilon^2}. \quad (41)$$

Note that $K_t^{(i)}$ depends on the perceived measurement-noise variance $R'_i(k_i)$, whereas the innovation volatility factor $\sqrt{\sigma_c^2 + \sigma_\epsilon^2}$ reflects the true signal-noise variance in $d\xi_{i,t}$.

Equivalently, since $d\xi_{i,t} - \hat{v}_{i,t} dt = \sqrt{\sigma_c^2 + \sigma_\epsilon^2} dI_{i,t}^\xi$, the Kalman–Bucy update (35) can be written in innovation form as

$$d\hat{v}_{i,t} = \mu_v dt + \beta_t^{(i)} dI_{i,t}^\xi.$$

This is the diffusion term that enters the reduced state dynamics (40).

(Under constant gain, $\beta_t^{(i)} \approx \beta_i^* = K_i^* \sqrt{\sigma_c^2 + \sigma_\epsilon^2}$ is constant.)

Diffusions entering the HJB. In the innovation representation (39)–(40), the controlled state is $(W_{i,t}, y_{i,t})$ and is driven by two independent Brownian motions: the *price innovation* I_i^p with volatility ϕ and the *filter innovation* I_i^ξ with volatility $\beta_t^{(i)}$. Consequently, the HJB/Itô generator contains the variance terms $\frac{1}{2} x_{i,t}^2 \phi^2 V_{ww}$ and $\frac{1}{2} ((\beta_t^{(i)})^2 + \phi^2) V_{yy}$. Moreover, because the same $dI_{i,t}^p$ appears in both $dW_{i,t}$ and $dy_{i,t}$, we have the cross-variation $d\langle W_i, y_i \rangle_t = -x_{i,t} \phi^2 dt$, which produces the mixed derivative term $-x_{i,t} \phi^2 V_{wy}$ in the HJB and is exactly the source of the intertemporal hedging correction in (48).

Control objective. Agent i maximizes exponential utility of terminal wealth:

$$\sup_{x_i \in \mathcal{A}} \mathbb{E} \left[-\exp(-\gamma_i W_{i,T}) \mid \mathcal{F}_t^i \right], \quad (42)$$

over admissible progressively measurable controls $\mathcal{A}$ with $\mathbb{E} \int_0^T x_{i,t}^2 dt < \infty$.

Exact CARA/LQG solution (finite horizon).

Proposition 4.3 (Exponential-quadratic value and optimal feedback). *Fix a mean-field trading-rate process $(\bar{u}_t)_{t \in [0, T]}$. The value function admits the exponential-quadratic form*

$$V_i(t, w, y) = -\exp \left(-\gamma_i w - \frac{1}{2} A_t^{(i)} y^2 - B_t^{(i)} y - C_t^{(i)} \right), \quad (43)$$

where $(A_t^{(i)}, B_t^{(i)}, C_t^{(i)})$ solve the backward ODE system with terminal conditions $A_T^{(i)} = B_T^{(i)} = C_T^{(i)} = 0$:

$$-\dot{A}_t^{(i)} = \frac{\kappa^2}{\phi^2} - (\beta_t^{(i)})^2 (A_t^{(i)})^2, \quad (44)$$

$$-\dot{B}_t^{(i)} = \frac{\kappa \lambda}{\phi^2} \bar{u}_t + \mu_v A_t^{(i)} - (\beta_t^{(i)})^2 A_t^{(i)} B_t^{(i)}, \quad (45)$$

$$-\dot{C}_t^{(i)} = \frac{1}{2} \left(((\beta_t^{(i)})^2 + \phi^2) A_t^{(i)} - (\beta_t^{(i)})^2 (B_t^{(i)})^2 \right) \quad (46)$$

$$+ \mu_v B_t^{(i)} + \frac{\lambda^2}{2\phi^2} \bar{u}_t^2. \quad (47)$$

The optimal control is the affine feedback

$$x_{i,t}^* = \frac{\kappa y_{i,t} + \lambda \bar{u}_t}{\gamma_i \phi^2} + \frac{A_t^{(i)} y_{i,t} + B_t^{(i)}}{\gamma_i}. \quad (48)$$

The first term is the myopic (instantaneous) demand. The second term is an intertemporal hedging correction induced by the correlation between the state $y_{i,t}$ and traded return noise in (40).

Constant-gain closed form (tight approximation). Under the constant-gain approximation $K_t^{(i)} \approx K_i^*$ so that $\beta_t^{(i)} \approx \beta_i^*$ is constant, (44) has the explicit solution

$$\begin{aligned} A_t^{(i)} &\approx \frac{\kappa}{\phi \beta_i^*} \tanh\left(\frac{\kappa \beta_i^*}{\phi}(T-t)\right), \\ \beta_i^* &:= K_i^* \sqrt{\sigma_c^2 + \sigma_\epsilon^2}, \\ K_i^* &:= \frac{\sigma_v}{\sqrt{\sigma_c^2 + \sigma_\epsilon^2/k_i}}. \end{aligned} \quad (49)$$

The remaining $(B_t^{(i)}, C_t^{(i)})$ follow from linear ODEs once $\bar{u}_t$ is specified.

Mean-field consistency. With flow-based impact, the endogenous mean-field object is the trading rate $\bar{u}_t$ rather than the inventory $\bar{x}_t$. Closing the loop therefore requires an explicit link between inventory adjustments and trading rates; Subsection 2.3.3 fixes the discrete-time timing used in the baseline particle implementation, and Subsection 4.4 sketches a continuous-time flow-control extension.

4.4 Continuous-time CARA/LQG with inventory dynamics and quadratic trading costs (flow-control extension)

This subsection provides a continuous-time benchmark consistent with the stock–flow separation (inventory x versus trading rate u). It is included only as an extension and is not used in the baseline discrete-time simulations.

Let $u_{i,t}$ denote the agent’s trading rate and let inventory evolve as

$$dx_{i,t} = u_{i,t} dt. \quad (50)$$

Under the innovation representation (Proposition 4.3), price evolves with flow-based impact,

$$dp_t = (\kappa y_{i,t} + \lambda \bar{u}_t) dt + \phi dI_{i,t}^p, \quad (51)$$

and we augment wealth with quadratic trading and holding costs,

$$dW_{i,t} = x_{i,t} dp_t - \eta u_{i,t}^2 dt - \frac{\chi}{2} x_{i,t}^2 dt, \quad \eta > 0, \chi > 0. \quad (52)$$

Consider the exponential-utility value function with state (w, x, y) ,

$$V(t, w, x, y) := \sup_{(u_s)_{s \in [t, T]}} \mathbb{E}[-\exp(-\gamma W_{i,T}) \mid (W_{i,t}, x_{i,t}, y_{i,t}) = (w, x, y)],$$

taking the mean-field path $\bar{u}_t$ as given. With the CARA ansatz

$$V(t, w, x, y) = -\exp(-\gamma(w + f(t, x, y))),$$

the HJB takes the standard risk-sensitive LQG form, and the control enters only through

$$\sup_u \{u \partial_x f(t, x, y) - \eta u^2\},$$

so the optimal trading rate is

$$u_{i,t}^* = \frac{\partial_x f(t, x_{i,t}, y_{i,t})}{2\eta}. \quad (53)$$

Under a quadratic specification for $f(t, x, y)$, $\partial_x f$ is affine in (x, y) and hence $u_{i,t}^*$ is linear in (x, y) . Substituting the quadratic ansatz into the HJB yields a coupled system of Riccati ODEs for the quadratic coefficients and linear ODEs for the affine terms, with forcing that depends on the (given) mean-field flow $\bar{u}_t$; this is the usual LQG structure.

4.5 Level Amplification vs. Overconfidence Effects

This subsection distinguishes two notions: (i) *level amplification*—large mispricing/volatility levels driven by weak anchoring and/or high noise trading; and (ii) *overconfidence (treatment) effects*—incremental differences when increasing k from 1 to 3. The following proposition gives an analytical mechanism for why k predominantly affects disagreement rather than mispricing.

Proposition 4.4 (Mechanism: mean-belief closure and why k shifts dispersion more than price). *Within the myopic mean-field equilibrium (per-step stock-flow closure in Proposition 4.5):*

- (i) **Closure depends only on the mean belief.** *The per-step fixed point for $(\bar{x}_t, \bar{u}_t)$ depends on the belief distribution only through cross-sectional averages of $\hat{v}_{i,t} - p_t$ (equivalently, the mean belief $m_t := \int_0^1 \hat{v}_{i,t} di$) and the predetermined pre-trade mean inventory $\bar{x}_t$. In particular, cross-sectional variance (disagreement) does not enter the mean-field fixed point except through the risk denominators $D_{i,t}$ (and in the homogeneous benchmark $D_{i,t}$ is common across agents). (In the particle implementation, the one-step myopic rule uses the effective variance $\text{Var}_{i,t}^{\text{eff}} = \phi^2 + \kappa^2 \Sigma_{i,t} \Delta t$ (Proposition 4.5); in the homogeneous benchmark $\Sigma_{i,t}$ is common across agents, so this term only rescales $\bar{x}_t$ and does not introduce dependence on disagreement.)*
- (ii) **Decomposition of mispricing.** *For additional closed-form intuition, under the reduced-form position-impact benchmark closure (31) (which collapses stock and flow), let $y_t := p_t - v_t$ (mispricing) and $e_t := m_t - v_t$ (mean-belief tracking error). Under the constant-gain approximation, the pair (y_t, e_t) satisfies the linear system (65): mispricing dynamics are driven by (a) microstructure regime (κ, λ, ϕ) through the coefficient $c := \lambda\kappa/(\gamma\phi^2 - \lambda)$ and the mean-reversion rate $\kappa + c$; and (b) the mean-belief term e_t , which is the only channel through which the belief distribution affects price. Thus $\mathbb{E}[|p - v|]$ and stationary mispricing variance are determined by (κ, λ, ϕ) and the stationary law of e_t .*
- (iii) **Why k predominantly affects disagreement.** *Under the information structure of Appendix A, each agent's belief has a common component (loading $K^*(k)\sigma_c$ on dU_t) and an idiosyncratic component (loading $K^*(k)\sigma_\epsilon$ on $dB_{i,t}$). The cross-sectional mean m_t aggregates beliefs, so the idiosyncratic components average out and m_t is driven only by the common noise and the fundamental. By Proposition 4.8, disagreement $D_\infty(k) = \text{Var}(\hat{v}_{i,t} - m_t \mid \mathcal{F}^0)$ is strictly increasing in k . By Proposition 4.9, the mean-belief error $e_t = m_t - v_t$ has stationary variance $\text{Var}(e_\infty)$ that depends on k only through $K^*(k)$ and is strictly decreasing in k . Hence incremental increases in k raise cross-sectional variance (disagreement) and trading intensity by construction, while the mean belief—and therefore equilibrium demand and mispricing—respond only weakly because they depend on e_t , whose variance is bounded and not uniformly amplified in k .*

Structural extensions (potential mechanisms). The model can be extended with structural mechanisms that could amplify mispricing persistence or volatility: (i) *State-dependent anchoring*: $\kappa(t) = \kappa_{\text{base}} \exp(-\alpha_\kappa |p_t - v_t|)$ weakens arbitrage pressure under extreme mispricing (limits to arbitrage); (ii) *Endogenous fundamental feedback*: $dv_t = \mu_v dt + \sigma_v dW_t + \beta_{\text{feedback}}(p_t - v_t) dt$ creates reflexive price-fundamental loops; (iii) *Overconfidence-dependent impact*: $\lambda_{\text{eff}} = \lambda_{\text{base}}(1 + \beta_\lambda(\bar{k}_t - 1))$ makes impact increase with aggregate overconfidence.

Proposition 4.5 (Per-step fixed point in the particle implementation). *In the discrete-time particle implementation, we use an effective variance*

$$\text{Var}_{i,t}^{\text{eff}} := \phi^2 + \kappa^2 \Sigma_{i,t} \Delta t, \quad (54)$$

i.e. we instantiate the variance rate in Proposition 4.2 as $\sigma_{p,i,t}^2 = \text{Var}_{i,t}^{\text{eff}}$, and define the effective risk/holding-cost term

$$D_{i,t} := \chi + \gamma \text{Var}_{i,t}^{\text{eff}}. \quad (55)$$

The myopic rule takes the linear form

$$x_{i,t} = \frac{\alpha_0}{D_{i,t}} \left(\kappa(\hat{v}_{i,t} - p_t) + \lambda \bar{u}_t \right). \quad (56)$$

This form follows by applying Proposition 4.2 to the flow-based Euler increment $\Delta p_{t+1} := p_{t+1} - p_t = \kappa(v_t - p_t)\Delta t + \lambda \bar{u}_t \Delta t + \phi \sqrt{\Delta t} \varepsilon_{t+1}^\eta$: conditional on $\mathcal{F}_t^i$, the drift depends on $\hat{v}_{i,t}$ and the variance satisfies $\text{Var}_{i,t}(\Delta p_{t+1}) = \phi^2 \Delta t + \kappa^2 \Sigma_{i,t} (\Delta t)^2 = \text{Var}_{i,t}^{\text{eff}} \Delta t$. Hence the optimal one-step inventory depends on the effective variance rate $\text{Var}_{i,t}^{\text{eff}} = \phi^2 + \kappa^2 \Sigma_{i,t} \Delta t$. The scalar $\alpha_0 > 0$ is an optional position-size (risk-budget) scaling (set to $\alpha_0 = 1$ in our baseline) and can be absorbed into a rescaling of γ . Define

$$A_t := \frac{1}{N} \sum_{i=1}^N \frac{\alpha_0 \kappa(\hat{v}_{i,t} - p_t)}{D_{i,t}}, \quad B_t := \frac{1}{N} \sum_{i=1}^N \frac{\alpha_0 \lambda}{D_{i,t}}.$$

Let $\bar{x}_{t-} := \frac{1}{N} \sum_{i=1}^N x_{i,t-}$ denote the pre-trade mean inventory and note that the mean trading rate satisfies $\bar{u}_t = (\bar{x}_t - \bar{x}_{t-})/\Delta t$. If $0 \leq B_t < \Delta t$, the scalar fixed point exists and is unique and is given by

$$\bar{x}_t = \frac{A_t - (B_t/\Delta t)\bar{x}_{t-}}{1 - B_t/\Delta t}, \quad \bar{u}_t = \frac{\bar{x}_t - \bar{x}_{t-}}{\Delta t}. \quad (57)$$

Theorem 4.6 (Contraction and uniqueness of the per-step mean-field closure). *Fix time t and treat the belief states $(\hat{v}_{i,t}, \Sigma_{i,t})$, the public price p_t , and the pre-trade mean inventory $\bar{x}_{t-}$ as given. Define the affine mean-field map*

$$\Phi_t(\bar{x}) := A_t + \frac{B_t}{\Delta t} \bar{x} - \frac{B_t}{\Delta t} \bar{x}_{t-}, \quad (58)$$

where A_t, B_t are as in Proposition 4.5. If $0 \leq B_t < \Delta t$, then Φ_t is a contraction on $\mathbb{R}$ with constant $B_t/\Delta t$ and has a unique fixed point, given in closed form in Proposition 4.5. Moreover, since $\text{Var}_{i,t}^{\text{eff}} \geq \phi^2$ and $D_{i,t} \geq \chi + \gamma \phi^2$, a simple sufficient condition ensuring $B_t < \Delta t$ uniformly is

$$\frac{\alpha_0 \lambda}{\Delta t (\chi + \gamma \phi^2)} < 1. \quad (59)$$

In the homogeneous case $D_{i,t} \equiv D_t$, we have $B_t = \alpha_0 \lambda / D_t$, so $0 \leq B_t < \Delta t$ is equivalent to $D_t > \alpha_0 \lambda / \Delta t$, i.e. $\chi + \gamma \text{Var}_t^{\text{eff}} > \alpha_0 \lambda / \Delta t$.

Proposition 4.7 (Closed-loop stability in a homogeneous benchmark). *Assume the homogeneous closure (31) holds and, for this benchmark calculation, that the mean belief tracks the fundamental ($m_t = v_t$). Then the mispricing $y_t := p_t - v_t$ satisfies the linear SDE*

$$\begin{aligned} dy_t &= -\kappa_{\text{eff}} y_t dt - \mu_v dt + \phi dZ_t - \sigma_v dW_t, \\ \kappa_{\text{eff}} &:= \kappa \frac{\gamma \phi^2}{\gamma \phi^2 - \lambda}. \end{aligned} \tag{60}$$

In particular, if $\gamma \phi^2 > \lambda$ then $\kappa_{\text{eff}} > 0$ and y_t is mean-reverting (Ornstein–Uhlenbeck). If additionally $\mu_v = 0$, the stationary variance is $\text{Var}(y_\infty) = (\phi^2 + \sigma_v^2)/(2\kappa_{\text{eff}})$.

4.6 Uniqueness and Stability Conditions (Multiplicity and Blow-Ups)

The following conditions rule out multiple equilibria and blow-ups. **Continuous-time LQG map Φ** (position-impact benchmark; Section 3.7): if $\gamma \phi^2 > \lambda$ (positive denominator) and $q < 1$, then Φ is a contraction on $(\mathcal{X}, \|\cdot\|_2)$, so the fixed point is unique and Picard iteration does not blow up (see Proposition 3.7). **Discrete-time per-step (baseline stock–flow)**: if $0 \leq B_t < \Delta t$, the scalar stock–flow fixed point for $(\bar{x}_t, \bar{u}_t)$ is unique (Proposition 4.5); a sufficient uniform condition is (59) (see Theorem 4.6). **Closed-loop paths (mispricing OU benchmark)**: under the position-impact benchmark closure (31), if $\gamma \phi^2 > \lambda$ then $\kappa_{\text{eff}} > 0$ and mispricing is mean-reverting (Proposition 4.7). **Black-box MFG limit (benchmark)**: under (H1)–(H5), the equilibrium flow is unique and $\bar{\mu}_t^N$ converges to μ_t at rate $O(1/N)$ (see Proposition 3.5 in the mean-field limit subsection).

4.7 Option B: Stationary Moments With Explicit k -Dependence

This subsection is a diagnostic benchmark built on the reduced-form position-impact closure (31); the baseline simulations use the discrete-time stock–flow equilibrium with uniqueness condition $0 \leq B_t < \Delta t$ (Proposition 4.5, Theorem 4.6).

Proposition 4.7 provides a transparent OU benchmark for mispricing, but it removes the overconfidence channel by imposing $m_t = v_t$. We now record complementary stationary-moment calculations in which k enters explicitly through the steady-state filtering gain (Appendix A). Lemma A.1 shows that the perceived Riccati equation converges exponentially to the steady state, motivating the constant-gain approximation used below (interpretable as an asymptotic stationary regime after burn-in).

Proposition 4.8 (Stationary disagreement and monotonicity in k (constant-gain approximation)). *Assume a homogeneous overconfidence type k and approximate the Kalman–Bucy gain by its steady-state value $K^*(k) = \sigma_v / \sqrt{R'(k)}$, where $R'(k) = \sigma_c^2 + \sigma_\epsilon^2/k$ (Appendix A). Let $m_t := \int_0^1 \hat{v}_{i,t} di$ and define the cross-sectional deviation $\delta_{i,t} := \hat{v}_{i,t} - m_t$. Conditioning on the common filtration $\mathcal{F}_t^0$, the deviations satisfy the OU dynamics*

$$d\delta_{i,t} = -K^*(k) \delta_{i,t} dt + K^*(k) \sigma_\epsilon dB_{i,t}. \tag{61}$$

Consequently, the conditional disagreement $D_t := \text{Var}(\hat{v}_{i,t} \mid \mathcal{F}_t^0) = \text{Var}(\delta_{i,t} \mid \mathcal{F}_t^0)$ admits the stationary value

$$D_\infty(k) = \frac{K^*(k) \sigma_\epsilon^2}{2} = \frac{\sigma_v \sigma_\epsilon^2}{2\sqrt{\sigma_c^2 + \sigma_\epsilon^2/k}}, \tag{62}$$

which is strictly increasing in k .

Proposition 4.9 (Stationary mean-belief tracking error (constant-gain approximation)). *Under the same constant-gain approximation, define the mean-belief tracking error $e_t := m_t - v_t$. Then*

$$de_t = -K^*(k) e_t dt + K^*(k) \sigma_c dU_t - \sigma_v dW_t, \quad (63)$$

and the stationary variance is

$$\text{Var}(e_\infty) = \frac{(K^*(k))^2 \sigma_c^2 + \sigma_v^2}{2K^*(k)} = \frac{K^*(k) \sigma_c^2}{2} + \frac{\sigma_v^2}{2K^*(k)}. \quad (64)$$

In particular, k enters the stationary tracking quality through $K^(k)$ and $\text{Var}(e_\infty)$ is strictly decreasing in k (and converges to $\sigma_v \sigma_c$ as $k \rightarrow \infty$ when $\sigma_c > 0$).*

Proposition 4.10 (Stationary mispricing with belief feedback (2D OU / Lyapunov form)). *Assume $\mu_v = 0$ and use the homogeneous closure (31). Under the constant-gain approximation above, the pair (y_t, e_t) with $y_t := p_t - v_t$ satisfies the linear system*

$$\begin{aligned} dy_t &= -(\kappa + c) y_t dt + c e_t dt + \phi dZ_t - \sigma_v dW_t, \\ de_t &= -K^*(k) e_t dt + K^*(k) \sigma_c dU_t - \sigma_v dW_t, \end{aligned} \quad (65)$$

where $c := \lambda\kappa/(\gamma\phi^2 - \lambda)$. If $\gamma\phi^2 > \lambda$ and $K^(k) > 0$, then the drift matrix*

$$A(k) := \begin{bmatrix} -(\kappa + c) & c \\ 0 & -K^*(k) \end{bmatrix}$$

is Hurwitz and the stationary covariance $P(k) = \mathbb{E}[X_\infty X_\infty^\top]$ for $X_t = (y_t, e_t)$ is the unique solution of the Lyapunov equation

$$A(k)P(k) + P(k)A(k)^\top + Q(k) = 0, \quad (66)$$

where $Q(k)$ is the diffusion covariance implied by (65). Thus $\text{Var}(y_\infty)$ depends on k explicitly through $K^(k)$.*

Corollary 4.11 (Closed-form stationary variances for Proposition 4.10). *In the setting of Proposition 4.10, write $K := K^*(k)$ and define the noise covariances $q_{11} := \phi^2 + \sigma_v^2$, $q_{12} := \sigma_v^2$, and $q_{22} := K^2 \sigma_c^2 + \sigma_v^2$. Then the Lyapunov equation (66) has the explicit solution $P = \begin{pmatrix} p_{11} & p_{12} \\ p_{12} & p_{22} \end{pmatrix}$ with*

$$\begin{aligned} p_{22} &= \frac{q_{22}}{2K} = \frac{K\sigma_c^2}{2} + \frac{\sigma_v^2}{2K}, \\ p_{12} &= \frac{c p_{22} + q_{12}}{\kappa + c + K}, \\ p_{11} &= \frac{c p_{12} + q_{11}/2}{\kappa + c}. \end{aligned} \quad (67)$$

In particular, $\text{Var}(y_\infty) = p_{11}$ depends on k through $K^(k)$.*

4.8 Connection to Numerical Results

Section VI simulates a large- N particle system implementing (8), Kalman belief updates (Appendix A), and the feedback policy (29) (implemented in discrete time with an effective variance that includes belief uncertainty). Overconfidence enters only through belief updating (misperceived signal precision), and we report price/mispricing/volatility statistics across bull/bear drifts μ_v and overconfidence levels k .

5 PROOFS

Proof of Proposition 4.1

Proof. From the static price equation (5), the coefficient on v is

$$a(k) = \frac{\lambda \tau'_{\epsilon,0}}{\gamma + \lambda(\tau_{v,0} + \tau'_{\epsilon,0})}, \quad \tau'_{\epsilon,0} = \frac{k}{\sigma_{\epsilon,0}^2}, \quad \tau_{v,0} = \frac{1}{\sigma_{v,0}^2}.$$

Set $u := \tau'_{\epsilon,0} > 0$ and define $f(u) := \frac{\lambda u}{\gamma + \lambda(\tau_{v,0} + u)}$. Then $a(k) = f(\tau'_{\epsilon,0})$ and

$$f'(u) = \frac{\lambda(\gamma + \lambda\tau_{v,0})}{(\gamma + \lambda(\tau_{v,0} + u))^2} > 0,$$

so f is strictly increasing in u , hence $a(k)$ is strictly increasing in k . Moreover, $a(k) > 0$ is immediate since all parameters are positive, and

$$a(k) < 1 \iff \lambda u < \gamma + \lambda(\tau_{v,0} + u) \iff 0 < \gamma + \lambda\tau_{v,0},$$

which holds. Thus $0 < a(k) < 1$ and $a(k)$ increases with k . $\square$

Proof of Proposition 4.2

Proof. Fix a decision interval Δt and condition on $\mathcal{F}_t^i$. By the proposition's assumptions, $\Delta p_{t+\Delta t}$ is Gaussian with

$$\mu_{i,t} := \mathbb{E}_i[\Delta p_{t+\Delta t}] = (\kappa(\hat{v}_{i,t} - p_t) + \lambda \bar{u}_t) \Delta t, \quad \text{Var}_i(\Delta p_{t+\Delta t}) = \sigma_{p,i,t}^2 \Delta t.$$

The one-step wealth increment is

$$\Delta W_{i,t} = x_{i,t} \Delta p_{t+\Delta t} - \frac{\chi}{2} x_{i,t}^2 \Delta t,$$

hence $\Delta W_{i,t}$ is Gaussian conditional on $\mathcal{F}_t^i$ with

$$\mathbb{E}_i[\Delta W_{i,t}] = x_{i,t} \mu_{i,t} - \frac{\chi}{2} x_{i,t}^2 \Delta t, \quad \text{Var}_i(\Delta W_{i,t}) = x_{i,t}^2 \sigma_{p,i,t}^2 \Delta t.$$

For CARA utility $U(w) = -e^{-\gamma w}$ and Gaussian $X \sim N(m, s^2)$, the standard identity gives $\mathbb{E}[-e^{-\gamma X}] = -\exp(-\gamma m + (\gamma^2/2)s^2)$. Applying this with $X = W_{i,t} + \Delta W_{i,t}$ (where $W_{i,t}$ is $\mathcal{F}_t^i$ -measurable), maximizing conditional expected utility over $x_{i,t}$ is equivalent to maximizing the certainty-equivalent quadratic

$$x_{i,t} \mu_{i,t} - \frac{\gamma}{2} x_{i,t}^2 \sigma_{p,i,t}^2 \Delta t - \frac{\chi}{2} x_{i,t}^2 \Delta t.$$

The first-order condition yields $\mu_{i,t} - (\gamma \sigma_{p,i,t}^2 + \chi) x_{i,t} \Delta t = 0$, so

$$x_{i,t}^* = \frac{\mu_{i,t} / \Delta t}{\chi + \gamma \sigma_{p,i,t}^2} = \frac{\kappa(\hat{v}_{i,t} - p_t) + \lambda \bar{u}_t}{\chi + \gamma \sigma_{p,i,t}^2},$$

which is (29). $\square$

Proof of Proposition 4.5

Proof. Work on the discrete grid with step Δt and the Euler increment

$$\begin{aligned}\Delta p_{t+1} &:= p_{t+1} - p_t = \kappa(v_t - p_t)\Delta t + \lambda \bar{u}_t \Delta t \\ &\quad + \phi \sqrt{\Delta t} \varepsilon_{t+1}^\eta, \\ \varepsilon_{t+1}^\eta &\sim N(0, 1).\end{aligned}$$

Conditional on investor i 's information $\mathcal{F}_t^i$, we have $\mathbb{E}[v_t \mid \mathcal{F}_t^i] = \hat{v}_{i,t}$ and $\text{Var}(v_t \mid \mathcal{F}_t^i) = \Sigma_{i,t}$. Hence

$$\mathbb{E}[\Delta p_{t+1} \mid \mathcal{F}_t^i] = (\kappa(\hat{v}_{i,t} - p_t) + \lambda \bar{u}_t) \Delta t,$$

and, using independence of ε_{t+1}^η from $\mathcal{F}_t^i$ and from v_t ,

$$\begin{aligned}\text{Var}(\Delta p_{t+1} \mid \mathcal{F}_t^i) &= \kappa^2 \text{Var}(v_t \mid \mathcal{F}_t^i) (\Delta t)^2 + \phi^2 \Delta t \\ &= \kappa^2 \Sigma_{i,t} (\Delta t)^2 + \phi^2 \Delta t.\end{aligned}$$

Under the one-step myopic CARA approximation (Proposition 4.2), the investor chooses $x_{i,t}$ to maximize the certainty-equivalent quadratic

$$x_{i,t} \mathbb{E}[\Delta p_{t+1} \mid \mathcal{F}_t^i] - \frac{\gamma}{2} x_{i,t}^2 \text{Var}(\Delta p_{t+1} \mid \mathcal{F}_t^i) - \frac{\chi}{2} x_{i,t}^2 \Delta t.$$

Substituting the conditional moments above and dividing by the common factor $\Delta t > 0$ gives the equivalent problem

$$\max_{x_{i,t}} x_{i,t} (\kappa(\hat{v}_{i,t} - p_t) + \lambda \bar{u}_t) - \frac{1}{2} \left(\gamma(\phi^2 + \kappa^2 \Sigma_{i,t} \Delta t) + \chi \right) x_{i,t}^2.$$

The unique maximizer is

$$x_{i,t}^* = \frac{\kappa(\hat{v}_{i,t} - p_t) + \lambda \bar{u}_t}{\chi + \gamma(\phi^2 + \kappa^2 \Sigma_{i,t} \Delta t)}.$$

Including the optional scale $\alpha_0 > 0$ (risk-budget/position-size parameter) yields (56) with

$$\text{Var}_{i,t}^{\text{eff}} := \phi^2 + \kappa^2 \Sigma_{i,t} \Delta t,$$

Write (56) as $x_{i,t} = a_{i,t} + b_{i,t} \bar{u}_t$ with

$$a_{i,t} := \frac{\alpha_0 \kappa(\hat{v}_{i,t} - p_t)}{\chi + \gamma \text{Var}_{i,t}^{\text{eff}}}, \quad b_{i,t} := \frac{\alpha_0 \lambda}{\chi + \gamma \text{Var}_{i,t}^{\text{eff}}}.$$

Averaging over $i = 1, \dots, N$ gives $\bar{x}_t = A_t + B_t \bar{u}_t$ where A_t, B_t are as defined in the paper. Using $\bar{u}_t = (\bar{x}_t - \bar{x}_{t-})/\Delta t$ gives

$$\bar{x}_t = A_t + \frac{B_t}{\Delta t} \bar{x}_t - \frac{B_t}{\Delta t} \bar{x}_{t-}.$$

Rearranging yields $(1 - B_t/\Delta t) \bar{x}_t = A_t - (B_t/\Delta t) \bar{x}_{t-}$. If $0 \leq B_t < \Delta t$, the unique solution is

$$\bar{x}_t = \frac{A_t - (B_t/\Delta t) \bar{x}_{t-}}{1 - B_t/\Delta t}, \quad \bar{u}_t = \frac{\bar{x}_t - \bar{x}_{t-}}{\Delta t},$$

as claimed. $\square$

Proof of Theorem 4.6

Proof. By Proposition 4.5, for fixed $(\hat{v}_{i,t}, \Sigma_{i,t})_{i=1}^N$, p_t , and $\bar{x}_{t-}$, the per-step map is affine:

$$\Phi_t(\bar{x}) = A_t + \frac{B_t}{\Delta t} \bar{x} - \frac{B_t}{\Delta t} \bar{x}_{t-}.$$

For any $\bar{x}, \bar{x}'$,

$$|\Phi_t(\bar{x}) - \Phi_t(\bar{x}')| = \frac{|B_t|}{\Delta t} |\bar{x} - \bar{x}'|.$$

Thus if $0 \leq B_t < \Delta t$, Φ_t is a contraction on $\mathbb{R}$ with contraction constant $B_t/\Delta t$ and has a unique fixed point, given in closed form in Proposition 4.5.

For the sufficient uniform condition: since $\Sigma_{i,t} \geq 0$ and $\Delta t > 0$, we have $\text{Var}_{i,t}^{\text{eff}} = \phi^2 + \kappa^2 \Sigma_{i,t} \Delta t \geq \phi^2$, hence $D_{i,t} = \chi + \gamma \text{Var}_{i,t}^{\text{eff}} \geq \chi + \gamma \phi^2$ and therefore

$$B_t = \frac{1}{N} \sum_{i=1}^N \frac{\alpha_0 \lambda}{D_{i,t}} \leq \frac{1}{N} \sum_{i=1}^N \frac{\alpha_0 \lambda}{\chi + \gamma \phi^2} = \frac{\alpha_0 \lambda}{\chi + \gamma \phi^2}.$$

Therefore $\alpha_0 \lambda / (\Delta t (\chi + \gamma \phi^2)) < 1$ implies $B_t < \Delta t$ for all t . $\square$

Proof of Theorem 3.1 (Well-posedness of the discrete-time particle system)

Proof. Fix $T \in \mathbb{N}$, $\Delta t > 0$, and fix a realization of all common and idiosyncratic shock sequences used by the algorithm. We show by induction on $t = 0, 1, \dots, T$ that the state

$$S_t = (v_t, p_t, \{(\hat{v}_{i,t}, \Sigma_{i,t})\}_{i=1}^N, \bar{x}_t)$$

is uniquely determined and adapted.

At $t = 0$, $(v_0, p_0, \{\hat{v}_{i,0}, \Sigma_{i,0}\}_{i=1}^N)$ are given by assumption (A1). Suppose the system is uniquely defined up to time t . Then:

1. By (A2), the discrete-time Kalman recursion in Appendix A defines each $(\hat{v}_{i,t}, \Sigma_{i,t})$ uniquely from $(\hat{v}_{i,t-1}, \Sigma_{i,t-1})$ and the realized private signal at time t (which is a measurable function of the fixed shock sequences and v_t).
2. Given $(\hat{v}_{i,t}, \Sigma_{i,t})_{i=1}^N$ and p_t , the per-step demand map Φ_t is well-defined by Proposition 4.5. By (A3) (e.g. $B_t < \Delta t$), Φ_t has a unique fixed point $\bar{x}_t$ (and hence $\bar{u}_t$), and then each $x_{i,t}$ is uniquely determined.
3. Finally, by (A4), the Euler updates for v_{t+1} and p_{t+1} are explicit measurable functions of (v_t, p_t) , $\bar{u}_t$, and the realized shocks at time $t + 1$, hence define unique (v_{t+1}, p_{t+1}) .

Thus S_{t+1} is uniquely determined from S_t and the fixed shocks, and is adapted to the natural filtration generated by the shocks. Induction completes the proof, yielding a unique adapted solution path on $\{0, 1, \dots, T\}$. $\square$

Proof of Proposition 4.8 (Constant-gain approximation)

Proof. Assume a homogeneous type k and adopt the constant-gain approximation $K_t \equiv K^*(k) =: K > 0$. From Appendix A, the (subjective) Kalman–Bucy filter is

$$d\hat{v}_{i,t} = \mu_v dt + K(d\xi_{i,t} - \hat{v}_{i,t} dt).$$

Using the observation equation (33),

$$d\xi_{i,t} = v_t dt + \sigma_c dU_t + \sigma_\epsilon dB_{i,t},$$

so substituting gives

$$\begin{aligned} d\hat{v}_{i,t} &= \mu_v dt + K((v_t - \hat{v}_{i,t}) dt + \sigma_c dU_t + \sigma_\epsilon dB_{i,t}) \\ &= (\mu_v + K(v_t - \hat{v}_{i,t})) dt + K\sigma_c dU_t + K\sigma_\epsilon dB_{i,t}. \end{aligned}$$

Let $m_t := \int_0^1 \hat{v}_{i,t} di$ be the cross-sectional mean belief. Under a continuum (or large- N) approximation, the idiosyncratic noises average out, i.e. $\int_0^1 dB_{i,t} di = 0$ in the sense of laws/LLN, while the common term U_t persists. Averaging the SDE over i yields

$$dm_t = (\mu_v + K(v_t - m_t)) dt + K\sigma_c dU_t.$$

Define the deviation $\delta_{i,t} := \hat{v}_{i,t} - m_t$. Subtracting the m_t equation from the $\hat{v}_{i,t}$ equation cancels the common-noise terms and gives

$$\begin{aligned} d\delta_{i,t} &= -K(\hat{v}_{i,t} - m_t) dt + K\sigma_\epsilon dB_{i,t} \\ &= -K\delta_{i,t} dt + K\sigma_\epsilon dB_{i,t}, \end{aligned}$$

which is (61).

Now condition on the common filtration $\mathcal{F}_t^0 := \sigma(W_s, U_s, Z_s : s \leq t)$. Conditional on $\mathcal{F}_t^0$, $\delta_{i,t}$ solves an OU SDE driven only by the idiosyncratic Brownian motion $B_{i,t}$, so its conditional variance $D_t := \text{Var}(\delta_{i,t} | \mathcal{F}_t^0) = \mathbb{E}[\delta_{i,t}^2 | \mathcal{F}_t^0]$ satisfies (Itô + tower property)

$$\begin{aligned} d(\delta_{i,t}^2) &= 2\delta_{i,t} d\delta_{i,t} + (d\delta_{i,t})^2 \\ &= (-2K\delta_{i,t}^2 + K^2\sigma_\epsilon^2) dt + 2K\sigma_\epsilon\delta_{i,t} dB_{i,t}. \end{aligned}$$

Taking $\mathcal{F}_t^0$ -conditional expectation kills the martingale term and yields the linear ODE

$$\frac{d}{dt} D_t = -2KD_t + K^2\sigma_\epsilon^2.$$

Hence $D_t \rightarrow D_\infty := \frac{K^2\sigma_\epsilon^2}{2K} = \frac{K\sigma_\epsilon^2}{2}$, giving (62).

Finally, with $R'(k) = \sigma_c^2 + \sigma_\epsilon^2/k$ we have $K^*(k) = \sigma_v / \sqrt{R'(k)}$, which is strictly increasing in k because $R'(k)$ is strictly decreasing in k ; therefore $D_\infty(k) = (K^*(k)\sigma_\epsilon^2)/2$ is strictly increasing in k . $\square$

Proof of Proposition 4.9 (Constant-gain approximation)

Proof. Under the constant-gain approximation (same setting as Proposition 4.8), we derived

$$dm_t = (\mu_v + K(v_t - m_t)) dt + K\sigma_c dU_t.$$

The fundamental satisfies $dv_t = \mu_v dt + \sigma_v dW_t$. Define $e_t := m_t - v_t$. Then

$$\begin{aligned} de_t &= dm_t - dv_t \\ &= (\mu_v + K(v_t - m_t)) dt + K\sigma_c dU_t - \mu_v dt - \sigma_v dW_t \\ &= -Ke_t dt + K\sigma_c dU_t - \sigma_v dW_t, \end{aligned}$$

which is (63).

Since U and W are independent Brownian motions, e_t is an OU process with mean-reversion rate K and diffusion variance rate $K^2\sigma_c^2 + \sigma_v^2$. The stationary variance is the standard OU value

$$\text{Var}(e_\infty) = \frac{K^2\sigma_c^2 + \sigma_v^2}{2K} = \frac{K\sigma_c^2}{2} + \frac{\sigma_v^2}{2K},$$

which is (64).

To see monotonicity in k , view $\text{Var}(e_\infty)$ as a function of $K > 0$:

$$g(K) := \frac{K\sigma_c^2}{2} + \frac{\sigma_v^2}{2K}, \quad g'(K) = \frac{\sigma_c^2}{2} - \frac{\sigma_v^2}{2K^2}.$$

If $\sigma_c = 0$, then $g(K) = \sigma_v^2/(2K)$ is strictly decreasing in K , and since $K^*(k)$ is strictly increasing in k , $\text{Var}(e_\infty)$ is strictly decreasing in k . If $\sigma_c > 0$, then $g'(K) < 0$ on $K \in (0, \sigma_v/\sigma_c)$ and $K^*(k) = \sigma_v/\sqrt{R'(k)} \leq \sigma_v/\sigma_c$ since $R'(k) = \sigma_c^2 + \sigma_\epsilon^2/k \geq \sigma_c^2$. Thus g is decreasing along the range of $K^*(k)$, so $\text{Var}(e_\infty)$ is strictly decreasing in k . $\square$

Proof of Proposition 4.10 (2D OU / Lyapunov form)

Proof. Assume $\mu_v = 0$ and the homogeneous closure (31):

$$\bar{x}_t = \frac{\kappa}{\gamma\phi^2 - \lambda}(m_t - p_t).$$

Write $y_t := p_t - v_t$ and $e_t := m_t - v_t$. Using $m_t - p_t = (m_t - v_t) - (p_t - v_t) = e_t - y_t$, the price SDE (8) becomes

$$\begin{aligned} dp_t &= \kappa(v_t - p_t) dt + \lambda \bar{x}_t dt + \phi dZ_t \\ &= -\kappa y_t dt + \frac{\lambda\kappa}{\gamma\phi^2 - \lambda}(e_t - y_t) dt + \phi dZ_t. \end{aligned}$$

Let $c := \lambda\kappa/(\gamma\phi^2 - \lambda)$. Then

$$dp_t = -(\kappa + c)y_t dt + ce_t dt + \phi dZ_t.$$

Since $\mu_v = 0$, $dv_t = \sigma_v dW_t$, hence

$$dy_t = dp_t - dv_t = -(\kappa + c)y_t dt + ce_t dt + \phi dZ_t - \sigma_v dW_t.$$

From Proposition 4.9, the mean-belief tracking error satisfies

$$de_t = -Ke_t dt + K\sigma_c dU_t - \sigma_v dW_t, \quad K := K^*(k) > 0.$$

Stacking $X_t = (y_t, e_t)^\top$ yields the linear system (65) with drift matrix

$$A(k) = \begin{pmatrix} -(\kappa + c) & c \\ 0 & -K \end{pmatrix}$$

and diffusion matrix $G(k)$ as stated.

If $\gamma\phi^2 > \lambda$ then $c > 0$ and $\kappa + c > 0$; also $K > 0$ by definition. The eigenvalues of $A(k)$ are $-(\kappa + c)$ and $-K$, hence strictly negative: $A(k)$ is Hurwitz. For a Hurwitz A , the continuous-time Lyapunov equation

$$AP + PA^\top + GG^\top = 0$$

has a unique symmetric positive semidefinite solution P , which equals the stationary covariance $P = \mathbb{E}[X_\infty X_\infty^\top]$. This gives (66) and shows $\text{Var}(y_\infty)$ depends on k through $K^*(k)$. $\square$

Proof of Corollary 4.11 (Closed-form stationary variances)

Proof. Let $A = \begin{pmatrix} -(\kappa + c) & c \\ 0 & -K \end{pmatrix}$ and $Q := GG^\top$. Because dZ_t, dU_t, dW_t are independent, from (65) we have

$$Q = \begin{pmatrix} \phi^2 + \sigma_v^2 & \sigma_v^2 \\ \sigma_v^2 & K^2 \sigma_c^2 + \sigma_v^2 \end{pmatrix} =: \begin{pmatrix} q_{11} & q_{12} \\ q_{12} & q_{22} \end{pmatrix}.$$

Write $P = \begin{pmatrix} p_{11} & p_{12} \\ p_{12} & p_{22} \end{pmatrix}$. Compute

$$AP = \begin{pmatrix} -(\kappa + c)p_{11} + cp_{12} & -(\kappa + c)p_{12} + cp_{22} \\ -Kp_{12} & -Kp_{22} \end{pmatrix},$$

$$PA^\top = \begin{pmatrix} -(\kappa + c)p_{11} + cp_{12} & -Kp_{12} \\ -(\kappa + c)p_{12} + cp_{22} & -Kp_{22} \end{pmatrix}.$$

Thus the Lyapunov equation $AP + PA^\top + Q = 0$ yields the entrywise system:

$$\begin{aligned} (2, 2) : & -2Kp_{22} + q_{22} = 0, \\ (1, 2) : & -(\kappa + c + K)p_{12} + cp_{22} + q_{12} = 0, \\ (1, 1) : & -2(\kappa + c)p_{11} + 2cp_{12} + q_{11} = 0. \end{aligned}$$

Solving sequentially gives

$$p_{22} = \frac{q_{22}}{2K} = \frac{K^2 \sigma_c^2 + \sigma_v^2}{2K} = \frac{K \sigma_c^2}{2} + \frac{\sigma_v^2}{2K},$$

$$p_{12} = \frac{cp_{22} + q_{12}}{\kappa + c + K}, \quad p_{11} = \frac{cp_{12} + q_{11}/2}{\kappa + c},$$

which matches the closed-form expressions in Corollary 4.11. $\square$

Proof of Proposition 4.7 (Homogeneous closed-loop stability benchmark)

Proof. Assume the homogeneous closure (31) holds and, for this benchmark, $m_t = v_t$. Then

$$\bar{x}_t = \frac{\kappa}{\gamma\phi^2 - \lambda}(v_t - p_t),$$

and substituting into (8) gives

$$\begin{aligned} dp_t &= \kappa(v_t - p_t) dt + \lambda \frac{\kappa}{\gamma\phi^2 - \lambda}(v_t - p_t) dt + \phi dZ_t \\ &= -\kappa_{\text{eff}}(p_t - v_t) dt + \phi dZ_t, \end{aligned}$$

where

$$\kappa_{\text{eff}} := \kappa \left(1 + \frac{\lambda}{\gamma\phi^2 - \lambda} \right) = \kappa \frac{\gamma\phi^2}{\gamma\phi^2 - \lambda}.$$

Define $y_t := p_t - v_t$. Using $dv_t = \mu_v dt + \sigma_v dW_t$,

$$dy_t = dp_t - dv_t = -\kappa_{\text{eff}} y_t dt - \mu_v dt + \phi dZ_t - \sigma_v dW_t,$$

as claimed. If $\gamma\phi^2 > \lambda$ then $\kappa_{\text{eff}} > 0$ and y_t is an OU process. When $\mu_v = 0$, y_t has diffusion variance rate $\phi^2 + \sigma_v^2$ and mean-reversion rate κ_{eff} , so its stationary variance is

$$\text{Var}(y_\infty) = \frac{\phi^2 + \sigma_v^2}{2\kappa_{\text{eff}}}.$$

$\square$

Algorithm 2 Particle approximation (one scenario)

- 1: Fix parameters, scenario (μ_v, k) , horizon T , population size N , and seeds $m = 1, \dots, M$.
 - 2: **for** $m = 1$ to M **do**
 - 3: Simulate one realization of common shocks $(\Delta W_t, \Delta U_t, \Delta Z_t)_{t=0}^{T-1}$.
 - 4: Initialize (v_0, p_0) and agent beliefs $(\hat{v}_{i,0}, \Sigma_{i,0})$ for $i = 1, \dots, N$.
 - 5: Initialize inventories $x_{i,0-}$ and wealth $W_{i,0}$ for $i = 1, \dots, N$.
 - 6: **for** $t = 0$ to $T - 1$ **do**
 - 7: For each agent i , draw idiosyncratic signal noise $\Delta B_{i,t}$ and update $(\hat{v}_{i,t}, \Sigma_{i,t})$ via the (perceived) Kalman recursion.
 - 8: Compute $(\bar{x}_t, \bar{u}_t)$ from the per-step stock-flow fixed point using pre-trade inventories (Proposition 4.5).
 - 9: For each agent i , set $x_{i,t} = x_{i,t}^*$ from the myopic rule given $\bar{u}_t$ and compute trading rate $u_{i,t} = (x_{i,t} - x_{i,t-})/\Delta t$.
 - 10: Update (v_{t+1}, p_{t+1}) using the Euler discretization of (8) (flow-based impact).
 - 11: Update wealth $W_{i,t+1} = W_{i,t} + x_{i,t}\Delta p_{t+1} - (\chi/2)x_{i,t}^2\Delta t$ and set $x_{i,t+1-} = x_{i,t}$.
 - 12: **end for**
 - 13: Record seed-level time series (price, fundamental, returns, mispricing, demand moments).
 - 14: **end for**
 - 15: Report scenario figures/tables by averaging statistics across seeds (Monte Carlo over the common and idiosyncratic shocks).
-

6 Numerical Examples (technical extract)

Scheme (map form). Let $S_t = (v_t, p_t, \{(\hat{v}_{i,t}, \Sigma_{i,t}, x_{i,t-})\}_{i=1}^N)$. For fixed shocks (common and idiosyncratic) the particle system defines a measurable one-step map

$$S_{t+1} = \Psi_{\Delta t, N}(S_t; \text{shocks}), \quad (68)$$

obtained by composing: (a) the Kalman update (Appendix A), (b) the per-step mean-field fixed point (Theorem 4.6), and (c) the Euler updates for (v, p) . Under the contraction condition (59), the fixed-point step is stable and the full map is well-defined (Theorem 3.1).

6.1 Parameter identification and calibration strategy

While full empirical estimation is beyond the scope of this paper, we provide an identification argument showing which observable moments pin down the key structural parameters $(\kappa, \lambda, \phi, \sigma_\epsilon, \sigma_v)$ in principle. This clarifies what data would be needed for calibration and establishes the model's testability.

Identification via observable moments. Under the myopic mean-field equilibrium, the following moments are observable (or proxied) from price and volume data:

- (1) **Return volatility** $\text{std}(r_t)$: In the model, “return” is the price increment $r_t := \Delta p_t$ (value units). Empirically one often works with log returns $\Delta \log P_t$; for small moves, $\Delta \log P_t \approx \Delta P_t/P_t$, so a simple scaling can map model increments to empirical log-return units when needed. From (8), the price increment variance is

$$\text{Var}(\Delta p_{t+1} \mid \mathcal{F}_t) \approx \left(\kappa^2 \text{Var}(v_t - p_t \mid \mathcal{F}_t) + \lambda^2 \text{Var}(\bar{x}_t \mid \mathcal{F}_t) + \phi^2 \right) \Delta t + \text{cross terms},$$

Table 1: Model and numerical parameters.

Parameter	Description	Baseline Value	Units
<i>Static Benchmark (Sec. 2.2)</i>			
$\mu_{v,0}$	Prior mean of v	—	value
$\sigma_{v,0}$	Prior sd of v	—	value
$\sigma_{\epsilon,0}$	Signal-noise sd	—	value
<i>Dynamic Model (Sec. 2.3+)</i>			
γ	Risk aversion	1.0	—
χ	Inventory holding cost	0.05	value/(share ² ·step)
k	Overconfidence (precision scale)	1.0, 3.0	—
μ_v	Fundamental drift (bull/bear)	± 0.02	value/step
σ_v	Fundamental diffusion vol.	0.10	value/ $\sqrt{\text{step}}$
σ_c	Common signal component	0.10	value/ $\sqrt{\text{step}}$
σ_ϵ	Idiosyncratic signal loading	0.80	value/ $\sqrt{\text{step}}$
κ	Fundamental anchoring	0.005	1/step
λ	Impact strength	0.20	value/share
ϕ	Price/noise-trading volatility	0.50	value/ $\sqrt{\text{step}}$
<i>Numerical Parameters</i>			
N	Number of agents	200	—
T	Horizon (steps)	400	step
M	Random seeds (replications)	50	—
p_0	Initial price	0.0	value
v_0	Initial fundamental	0.0	value

In steady state, $\text{std}(r_t) = \text{std}(\Delta p_t)$ depends on ϕ (direct noise trading), λ (impact amplification), and κ (mean reversion strength). The ratio λ/κ can be identified from the relative contribution of demand-driven vs. fundamental-driven price variance.

- (2) **Mean-reversion speed:** The autocorrelation of price deviations from a moving average of fundamentals (or a proxy like a slow-moving average of prices) identifies κ . Specifically, under (8), the half-life of price deviations is approximately $(\ln 2)/\kappa$, so κ can be estimated from the decay rate of price-f fundamental autocorrelations.

$$h = \frac{\ln 2}{\kappa}, \quad \kappa = \frac{\ln 2}{h}. \quad (69)$$

- (3) **Impact coefficient λ :** The contemporaneous correlation between order flow (aggregate demand proxy) and price changes, normalized by the variance of order flow, identifies λ up to the mean-field scaling. Volume proxies (e.g., signed volume or turnover) can serve as demand proxies.

- (4) **Fundamental volatility σ_v :** The long-run variance of fundamentals (estimated from a slow-moving filter of prices or from fundamental proxies like earnings/dividends) identifies σ_v . Alternatively, σ_v can be identified from the steady-state variance of the Kalman filter innovation process.
- (5) **Signal noise σ_ϵ :** The cross-sectional dispersion in beliefs (proxied by analyst forecast dispersion or survey-based disagreement) relative to the fundamental variance identifies the signal-to-noise ratio σ_v/σ_ϵ . Given σ_v , this pins down σ_ϵ .

Identification chain (estimated vs. inferred vs. normalized). From price-only data, one can (a) estimate κ from a half-life diagnostic (item 2) and (b) calibrate the effective price/noise-trading volatility ϕ to match return volatility (item 1, after the unit mapping to log-return units if desired). Identifying λ itself requires an order-flow proxy (item 3); given λ , the corresponding order-flow noise scale is inferred mechanically as $\sigma_\eta = \phi/\lambda$ (Section II, microstructure sketch). In our numerical work, λ is treated as a normalization choice (Table 1), and we therefore interpret ϕ as the empirically pinned-down price-noise scale conditional on that normalization.

6.2 Parameter regimes and stress diagnostics

6.2.1 Targeted joint-move grid

6.3 Extensions and validation checks

A Kalman–Bucy Filter With Misperceived Signal Precision

This appendix summarizes the Kalman–Bucy filter used for belief updating in Section II.

A.1 State and Observation Equations

The (unobserved) fundamental evolves as a linear diffusion

$$dv_t = \mu_v dt + \sigma_v dW_t, \quad (70)$$

and investor i observes a private signal

$$d\xi_{i,t} = v_t dt + \sigma_c dU_t + \sigma_\epsilon dB_{i,t}, \quad (71)$$

where W , U , and B_i are independent Brownian motions and U is common across investors.

A.2 Filter and Riccati Equation

Let $\hat{v}_{i,t}$ and $\Sigma_{i,t}$ denote investor i 's Kalman–Bucy belief mean and posterior variance computed under the *perceived* measurement-noise variance $R'_i(k_i)$ below. When $k_i = 1$, $\hat{v}_{i,t}$ coincides with the true conditional mean $\mathbb{E}[v_t | \mathcal{F}_t^i]$ and $\Sigma_{i,t}$ with the true conditional variance. When $k_i \neq 1$ they are best interpreted as *subjective* filter states under the investor's misspecified noise model.

A.2.1 Belief-mean SDE and diffusion decomposition

The belief state is $(\hat{v}_{i,t}, \Sigma_{i,t})$. Only $\hat{v}_{i,t}$ is stochastic; $\Sigma_{i,t}$ follows a scalar Riccati ODE. Starting from the Kalman–Bucy form

$$\begin{aligned} d\hat{v}_{i,t} &= \mu_v dt + K_{i,t}(d\xi_{i,t} - \hat{v}_{i,t} dt), \\ K_{i,t} &= \frac{\Sigma_{i,t}}{R'_i}, \quad R'_i = \sigma_c^2 + \frac{\sigma_\epsilon^2}{k_i}. \end{aligned} \quad (72)$$

and using the signal SDE $d\xi_{i,t} = v_t dt + \sigma_c dU_t + \sigma_\epsilon dB_{i,t}$, we obtain the explicit diffusion decomposition

$$d\hat{v}_{i,t} = (\mu_v + K_{i,t}(v_t - \hat{v}_{i,t})) dt + K_{i,t}\sigma_c dU_t + K_{i,t}\sigma_\epsilon dB_{i,t}. \quad (73)$$

Thus the belief-mean diffusion has (i) a *common* component $K_{i,t}\sigma_c dU_t$ and (ii) an *idiosyncratic* component $K_{i,t}\sigma_\epsilon dB_{i,t}$. The perceived posterior variance satisfies the Riccati ODE

$$\dot{\Sigma}_{i,t} = \sigma_v^2 - \frac{\Sigma_{i,t}^2}{R'_i}, \quad (74)$$

and carries no diffusion term.

Here R'_i is the *perceived* measurement-noise variance. Overconfidence corresponds to $k_i > 1$, i.e., investors treat their signal as more precise than it truly is. The common component $\sigma_c dU_t$ prevents the law-of-large-numbers collapse of aggregate signal noise in the continuum limit and is the source of common-noise randomness in Section III.

The scalar Riccati equation admits an explicit solution and converges exponentially to its steady state:

Lemma A.1 (Riccati convergence to steady state). *Fix $R'_i > 0$ and $\sigma_v > 0$ and let Σ_t solve (74) with $\Sigma_0 \geq 0$. Then $\Sigma_t \rightarrow \Sigma^* := \sigma_v \sqrt{R'_i}$ as $t \rightarrow \infty$ and there exist constants $C, c > 0$ (depending on Σ_0, σ_v, R'_i) such that*

$$|\Sigma_t - \Sigma^*| \leq C e^{-ct}, \quad c = \frac{2\sigma_v}{\sqrt{R'_i}}. \quad (75)$$

Consequently, the gain $K_t = \Sigma_t/R'_i$ converges exponentially to $K^* = \sigma_v/\sqrt{R'_i}$.

Proof. Fix $R'_i > 0$ and $\sigma_v > 0$ and consider the scalar Riccati ODE (74):

$$\frac{d\Sigma_t}{dt} = \sigma_v^2 - \frac{\Sigma_t^2}{R'_i}.$$

Let $\Sigma^* := \sigma_v \sqrt{R'_i}$ (the positive equilibrium) and write $a := \sigma_v/\sqrt{R'_i}$. Set $u_t := \Sigma_t/\Sigma^*$. Then u_t solves

$$\dot{u}_t = a(1 - u_t^2), \quad u_0 = \Sigma_0/\Sigma^*.$$

Define the ratio $\rho_t := \frac{1-u_t}{1+u_t}$ (well-defined since $u_t \geq 0$). Differentiating gives

$$\dot{\rho}_t = -2a \rho_t,$$

hence $\rho_t = \rho e^{-2at}$ with $\rho := \rho_0 = \frac{1-u_0}{1+u_0} = \frac{\Sigma^* - \Sigma_0}{\Sigma^* + \Sigma_0}$. Solving for u_t yields the explicit solution

$$\Sigma_t = \Sigma^* \frac{1 - \rho e^{-2at}}{1 + \rho e^{-2at}}.$$

In all cases with $\Sigma_0 \geq 0$, the denominator satisfies $1 + \rho e^{-2at} > 0$ and $\Sigma_t \geq 0$ for all t , and $\Sigma_t \rightarrow \Sigma^*$ as $t \rightarrow \infty$ since $e^{-2at} \rightarrow 0$.

Moreover,

$$\Sigma_t - \Sigma^* = \Sigma^* \left(\frac{1 - \rho e^{-2at}}{1 + \rho e^{-2at}} - 1 \right) = -\frac{2\Sigma^* \rho e^{-2at}}{1 + \rho e^{-2at}},$$

so the convergence is exponential with rate constant $c = 2a = 2\sigma_v/\sqrt{R'_i}$. For the prefactor, note: if $\rho \geq 0$ (equivalently $\Sigma_0 \leq \Sigma^*$), then $1 + \rho e^{-2at} \geq 1$ and

$$|\Sigma_t - \Sigma^*| \leq 2\Sigma^* |\rho| e^{-2at}.$$

If $\rho < 0$ (equivalently $\Sigma_0 > \Sigma^*$), then $1 + \rho e^{-2at} \geq 1 + \rho = 1 - |\rho|$ and

$$|\Sigma_t - \Sigma^*| \leq \frac{2\Sigma^*|\rho|}{1 - |\rho|} e^{-2at}.$$

Thus the claimed bound holds with, e.g., $C := 2\Sigma^*|\rho|$ when $\Sigma_0 \leq \Sigma^*$ and $C := \frac{2\Sigma^*|\rho|}{1 - |\rho|}$ when $\Sigma_0 > \Sigma^*$ (both depending only on Σ_0, σ_v, R'_i). Since $K_t = \Sigma_t/R'_i$, it follows that $K_t \rightarrow K^* = \Sigma^*/R'_i = \sigma_v/\sqrt{R'_i}$ at the same exponential rate. $\square$

A.3 Steady State

In steady state, (74) gives

$$\Sigma_i^* = \sigma_v \sqrt{R'_i} = \sigma_v \sqrt{\sigma_c^2 + \frac{\sigma_\epsilon^2}{k_i}}, \quad (76)$$

and the corresponding steady-state gain is $K_i^* = \Sigma_i^*/R'_i = \sigma_v/\sqrt{R'_i}$. Thus larger k increases the Kalman gain and reduces perceived posterior variance, generating more aggressive belief updating and, through the trading rule, larger demand responses to perceived mispricing.

A.4 Discrete-Time Recursion Used in Simulations

Section VI implements a discrete-time Kalman recursion consistent with the continuous-time signal model (72)–(73). On a grid with step size $\Delta t > 0$, the simulated fundamental evolves as

$$v_{t+1} = v_t + \mu_v \Delta t + \sigma_v \sqrt{\Delta t} \varepsilon_{t+1}, \quad \varepsilon_{t+1} \sim \mathcal{N}(0, 1),$$

and investor i observes a discrete-time signal obtained from the *flow* observation $d\xi_{i,t}$ by taking increments and normalizing by Δt :

$$z_{i,t} := \frac{\xi_{i,t+\Delta t} - \xi_{i,t}}{\Delta t} = v_t + \frac{\sigma_c}{\sqrt{\Delta t}} \varepsilon_t^c + \frac{\sigma_\epsilon}{\sqrt{\Delta t}} \varepsilon_{i,t}^i, \quad \varepsilon_t^c, \varepsilon_{i,t}^i \sim \mathcal{N}(0, 1),$$

where ε_t^c is common across investors and $\varepsilon_{i,t}^i$ is idiosyncratic. (With $\Delta t = 1$, this reduces to $z_{i,t} = v_t + \sigma_c \varepsilon_t^c + \sigma_\epsilon \varepsilon_{i,t}^i$, matching the implementation.) Thus $\text{Var}(z_{i,t} | v_t) = (\sigma_c^2 + \sigma_\epsilon^2)/\Delta t$; write $R(\Delta t) := (\sigma_c^2 + \sigma_\epsilon^2)/\Delta t$ for the objective measurement variance. Overconfidence scales only the idiosyncratic component, so the perceived measurement variance for $z_{i,t}$ is

$$R'_i(\Delta t) = \frac{\sigma_c^2 + \sigma_\epsilon^2/k_i}{\Delta t}.$$

Let $(\hat{v}_{i,t}, \Sigma_{i,t})$ denote the posterior mean/variance after processing $z_{i,t}$. The discrete-time Kalman recursion is:

$$\begin{aligned} \Sigma_{i,t}^- &= \Sigma_{i,t-1} + \sigma_v^2 \Delta t, \\ K_{i,t} &= \frac{\Sigma_{i,t}^-}{\Sigma_{i,t}^- + R'_i(\Delta t)}, \\ \hat{v}_{i,t} &= \hat{v}_{i,t-1} + \mu_v \Delta t + K_{i,t} (z_{i,t} - (\hat{v}_{i,t-1} + \mu_v \Delta t)), \\ \Sigma_{i,t} &= (1 - K_{i,t}) \Sigma_{i,t}^-. \end{aligned}$$

Given the mean trading rate $\bar{u}_t$ (with $\bar{u}_t \Delta t = \Delta \bar{x}_t$ under the stock-flow timing), the discrete-time price update (Euler form of (8)) used in the code is

$$\begin{aligned} p_{t+1} &= p_t + \kappa(v_t - p_t) \Delta t + \lambda \bar{u}_t \Delta t + \phi \sqrt{\Delta t} \varepsilon_{t+1}^\eta, \\ \varepsilon_{t+1}^\eta &\sim \mathcal{N}(0, 1). \end{aligned}$$

A.5 Price-Augmented Filtering (Robustness Variant)

Because prices are public and informative under (8), a fully consistent informational equilibrium would incorporate p into the observation set and solve a joint filtering/fixed-point problem. As a robustness check aligned with the implemented Euler update above, we construct an *implicit* noisy observation of v_t from the realized price increment:

$$\tilde{v}_t^{(p)} := p_t + \frac{(p_{t+1} - p_t) - \lambda \bar{u}_t \Delta t}{\kappa \Delta t} = v_t + \frac{\phi}{\kappa \sqrt{\Delta t}} \varepsilon_{t+1}^\eta.$$

Conditional on $\mathcal{F}_t$, this yields a Gaussian observation of v_t with measurement variance $R_p = \phi^2/(\kappa^2 \Delta t)$ (scaled by a factor $\sigma_{p,\text{obs}}^2$ in the code to probe robustness). We then apply an additional Kalman update using $\tilde{v}_t^{(p)}$ as a measurement (after processing the private signal), holding the state prediction fixed.

References

- Brad M. Barber and Terrance Odean. Boys will be boys: Gender, overconfidence, and common stock investment. *The Quarterly Journal of Economics*, 116(1):261–292, 2001. doi: 10.1162/003355301556400.
- Pierre Cardaliaguet and Charles-Albert Lehalle. Mean field game of controls and an application to trade crowding. *Mathematics and Financial Economics*, 12(3):335–363, 2018. doi: 10.1007/s11579-018-0206-7.
- René Carmona and François Delarue. *Probabilistic Theory of Mean Field Games with Applications I: Mean Field FBSDEs, Control, and Games*, volume 83 of *Probability Theory and Stochastic Modelling*. Springer, Cham, 2018a. doi: 10.1007/978-3-319-58920-6.
- René Carmona and François Delarue. *Probabilistic Theory of Mean Field Games with Applications II: Mean Field Games with Common Noise and Master Equations*, volume 84 of *Probability Theory and Stochastic Modelling*. Springer, Cham, 2018b. doi: 10.1007/978-3-319-58973-1.
- Rene Carmona, Jean-Pierre Fouque, and Li-Hsien Sun. Mean field games and systemic risk. *Communications in Mathematical Sciences*, 13(4):911–933, 2015. doi: 10.4310/CMS.2015.v13.n4.a4.
- Kent Daniel, David Hirshleifer, and Avanidhar Subrahmanyam. Investor psychology and security market under- and overreactions. *The Journal of Finance*, 53(6):1839–1885, 1998. doi: 10.1111/0022-1082.00077.
- J. Bradford De Long, Andrei Shleifer, Lawrence H. Summers, and Robert J. Waldmann. Noise trader risk in financial markets. *Journal of Political Economy*, 98(4):703–738, 1990.
- Nicolas Fournier and Arnaud Guillin. On the rate of convergence in Wasserstein distance of the empirical measure. *Probability Theory and Related Fields*, 162(3):707–738, 2015. doi: 10.1007/s00440-014-0583-8.
- Simon Gervais and Terrance Odean. Learning to be overconfident. *The Review of Financial Studies*, 14(1):1–27, 2001. doi: 10.1093/rfs/14.1.1.
- Olivier Guéant, Jean-Michel Lasry, and Pierre-Louis Lions. Mean field games and applications. In *Paris-Princeton Lectures on Mathematical Finance 2010*, volume 2003 of *Lecture Notes in Mathematics*, pages 205–266. Springer, 2010. doi: 10.1007/978-3-642-14660-2_3.

- Harrison Hong and Jeremy C. Stein. A unified theory of underreaction, momentum trading, and overreaction in asset markets. *Journal of Finance*, 54(6):2143–2184, 1999. doi: 10.1111/0022-1082.00184.
- Minyi Huang, Roland P. Malhamé, and Peter E. Caines. Large population stochastic dynamic games: closed-loop McKean-Vlasov systems and the Nash certainty equivalence principle. *IEEE Transactions on Automatic Control*, 51(11):1730–1755, 2006. doi: 10.1109/TAC.2006.884922.
- Aimé Lachapelle, Jean-Michel Lasry, Charles-Albert Lehalle, and Pierre-Louis Lions. Efficiency of the price formation process in presence of high frequency participants: a mean field game analysis. *Mathematics and Financial Economics*, 10(3):223–262, 2016. doi: 10.1007/s11579-016-0161-0.
- Daniel Lacker. A general characterization of the mean field limit for stochastic differential games. *Probability Theory and Related Fields*, 165(3–4):581–648, 2016. doi: 10.1007/s00440-015-0641-9.
- Daniel Lacker. On the convergence of closed-loop Nash equilibria for the infinite horizon mean-field LQG game. *ESAIM: Control, Optimisation and Calculus of Variations*, 24(2):437–469, 2018. doi: 10.1051/cocv/2017011.
- Terrance Odean. Volume, volatility, price, and profit when all traders are above average. *The Journal of Finance*, 53(6):1887–1934, 1998. doi: 10.1111/0022-1082.00078.